

Python	
Basic	<p>What is Python? Enlist some of its benefits.</p> <p>Python is a high-level, interpreted, and general-purpose programming language. It is known for its simple and readable syntax, which makes it easy to learn and use.</p> <p>Some benefits of Python include:</p> <ul style="list-style-type: none"> - Easy to learn and read – clean syntax that resembles English. - Versatile and cross-platform – works on Windows, macOS, Linux, etc. - Large standard library and ecosystem – many built-in modules and third-party packages. - Supports multiple programming paradigms – procedural, object-oriented, and functional programming. - Strong community support – many tutorials, forums, and resources available.
	<p>Can you tell us if Python is object-oriented or functional programming?</p> <p>Python is a multi-paradigm programming language, which means it supports both object-oriented programming (OOP) and functional programming.</p> <ul style="list-style-type: none"> - Object-oriented programming: Python allows you to define classes, create objects, and use inheritance and polymorphism. - Functional programming: Python supports functions as first-class objects, higher-order functions, map, filter, reduce, and lambda expressions. <p>So, Python is both object-oriented and functional, depending on how you write your code.</p>

	<p>What rules govern local and global variables in Python?</p> <p>In Python, variable scope determines where a variable can be accessed:</p> <ol style="list-style-type: none"> 1. Local variables are defined inside a function and can only be accessed within that function. They are created when the function starts and destroyed when it ends. 2. Global variables are defined outside any function and can be accessed anywhere in the file. 3. To modify a global variable inside a function, you must use the <code>global</code> keyword. 	<pre>python x = 10 # global variable def my_function(): y = 5 # local variable global x x += 1 print("Inside function:", x, y) my_function() print("Outside function:", x)</pre> <p>Output:</p> <pre>bash Inside function: 11 5 Outside function: 11</pre>	
	<p>Can you tell us what is slicing in Python?</p> <p>Slicing in Python is a way to extract a portion of a sequence, like a list, string, or tuple, by specifying a start index, stop index, and optional step.</p>	<pre>python my_list = [0, 1, 2, 3, 4, 5] print(my_list[1:5:2]) # Output: [1, 3] print(my_list[:3]) # Output: [0, 1, 2] print(my_list[3:]) # Output: [3, 4, 5]</pre>	

	<p>What is namespace in Python?</p> <p>A namespace in Python is a container that holds names (identifiers) and maps them to objects. In other words, it is a system that keeps track of all the variable names and their corresponding objects in memory.</p> <p>Types of namespaces:</p> <ol style="list-style-type: none"> 1. Local namespace – inside a function. 2. Global namespace – at the module level. 3. Built-in namespace – Python's built-in functions and exceptions. 	<pre>python x = 10 # global namespace def my_func(): y = 5 # local namespace print(y) my_func() print(x)</pre>	Copy code
	<p>What is pass in Python?</p> <p>"pass" is a null statement that allows your code to run without errors even if the function or class has no content yet.</p>	<pre>python def my_function(): pass # To be implemented later class MyClass: pass # Empty class</pre>	Copy code
	<p>Can you explain what is unittest in Python?</p> <p>"unittest is Python's built-in framework for writing and running unit tests to ensure that individual parts of the code work correctly."</p>	<pre>python import unittest def add(a, b): return a + b class TestAdd(unittest.TestCase): def test_addition(self): self.assertEqual(add(2, 3), 5) if __name__ == "__main__": unittest.main()</pre>	Copy code

	<p>What are negative indexes in Python?</p> <p>"Negative indexes in Python let you access elements from the end of a sequence. -1 is the last element, -2 the second last, and so on, so you don't need the sequence length to access items from the end."</p>	<pre>python my_list = [10, 20, 30, 40, 50] print(my_list[-1]) # Kết quả: 50 (phản từ cuối) print(my_list[-2]) # Kết quả: 40 (phản từ kế cuối)</pre>	Copy code	
	<p>What are ODBC modules in Python?</p> <p>"ODBC modules in Python let programs connect to databases via the ODBC standard, run SQL queries, and manage data. Common ones are pyodbc and pypyodbc."</p>	<pre>python import pyodbc conn = pyodbc.connect("DRIVER={SQL Server};SERVER=localhost;DATABASE=TestDB;UID=user;PWD=password") cursor = conn.cursor() cursor.execute("SELECT * FROM Employees") for row in cursor.fetchall(): print(row) conn.close()</pre>	Copy code	
	<p>How will you send an email from a Python Script?</p> <p>"You can send emails in Python using smtplib to connect to an SMTP server and email to format the message."</p>	<pre>python import smtplib from email.mime.text import MIMEText # Email content msg = MIMEText("Hello, this is a test email from Python.") msg['Subject'] = "Test Email" msg['From'] = "your_email@example.com" msg['To'] = "recipient@example.com" # Connect to SMTP server and send with smtplib.SMTP('smtp.example.com', 587) as server: server.starttls() # Secure the connection server.login('your_email@example.com', 'password') server.send_message(msg) print("Email sent successfully!")</pre>	Copy code	

	<p>What is PEP 8 and its importance?</p> <p>"PEP 8 is Python's style guide for writing clean, readable code. It ensures consistency, improves collaboration, and makes code easier to maintain."</p>	
	<p>What are the key features of Python?</p> <p>Python has several key features:</p> <ol style="list-style-type: none"> 1. Easy to learn and readable – simple syntax, close to English. 2. Interpreted language – runs code line by line. 3. Dynamically typed – no need to declare variable types. 4. Cross-platform – runs on Windows, macOS, Linux, etc. 5. Extensive standard library – many modules for different tasks. 6. Object-oriented and functional – supports multiple programming paradigms. 7. Open-source – free to use with strong community support. 	
	<p>What does it mean to be dynamically typed in Python?</p> <p>"Dynamically typed means you do not need to say the type of a variable. Python finds it when the program runs, and you can change the type later."</p>	<pre>python x = 10 # x is an integer x = "Hello" # now x is a string</pre>
	<p>What is a scope in Python?</p> <p>"Scope in Python is where a variable can be used. Types: local (inside a function), enclosing (nested functions), global (whole file), built-in (Python's own names)."</p>	<pre>python x = 10 # global scope def outer(): y = 5 # enclosing scope def inner(): z = 2 # local scope print(x, y, z) inner() outer()</pre>

	<p>How can Python script be executable on Unix?</p> <p>To make a Python script executable on Unix:</p> <ol style="list-style-type: none"> 1. Add the shebang line at the top of the script to specify the Python interpreter: #!/usr/bin/env python3 2. Make the script executable using the chmod command: chmod +x script.py 3. Run the script directly from the terminal: .script.py <p>This allows the script to be executed like a standalone program.</p>	
	<p>What is Docstring in Python?</p> <p>"A docstring is a special string in Python used to explain a module, class, or function. It goes in triple quotes at the start and can be seen with <code>__doc__</code>."</p>	<pre>python Copy code def add(a, b): """This function returns the sum of two numbers.""" return a + b print(add.__doc__) Output: bash Copy code This function returns the sum of two numbers.</pre>
	<p>What is init in Python?</p> <p>"<code>__init__</code> is a special method in a class that runs when a new object is made. It sets up the object's attributes."</p>	<pre>python Copy code class Person: def __init__(self, name, age): self.name = name self.age = age p = Person("Alice", 25) print(p.name, p.age) # Output: Alice 25</pre>

	<p>What are lists and tuples in Python?</p> <p>In Python:</p> <ul style="list-style-type: none"> - List is an ordered, mutable collection of items. You can change, add, or remove elements. Lists are defined using square brackets []. - Tuple is an ordered, immutable collection of items. Once created, its elements cannot be changed. Tuples are defined using parentheses ().. 	<pre>python # List my_list = [1, 2, 3] my_list[0] = 10 # Lists are mutable my_list.append(4) # Tuple my_tuple = (1, 2, 3) # my_tuple[0] = 10 # This will raise an error, tuples are immutable print(my_list) # Output: [10, 2, 3, 4] print(my_tuple) # Output: (1, 2, 3)</pre>	Copy code
	<p>What is the difference between Arrays and lists in Python?</p> <p>Lists can hold different types and are flexible. Arrays hold one type, use less memory, and are faster for large numeric data.</p>	<pre>python # List my_list = [1, "hello", 3.5] # Array import array my_array = array.array('i', [1, 2, 3]) # 'i' means integer type</pre>	Copy code
	<p>What is Self-used for in Python?</p> <p>"Self is a reference to the current object in a class. It lets you access the object's attributes and methods inside the class."</p>	<pre>python class Person: def __init__(self, name): self.name = name def greet(self): print("Hello, my name is", self.name) p = Person("Alice") p.greet() # Output: Hello, my name is Alice</pre>	Copy code

	<p>What are the major two-loop statements in Python?</p> <p>"Python has two main loops: for loops to go through items in a sequence, and while loops to repeat code while a condition is true."</p>	<pre>python # for loop for i in range(3): print(i) # while loop count = 0 while count < 3: print(count) count += 1</pre>	Copy code
	<p>What are Decorators in Python?</p> <p>A decorator in Python is a function that modifies or enhances another function or method without changing its code. Decorators are often used for logging, access control, timing, or caching. They are applied using the <code>@decorator_name</code> syntax above a function definition.</p>	<pre>❶ main.py > ... ❷ def my_decorator(func): ❸ def wrapper(): ❹ print("Before function call") ❺ func() ❻ print("After function call") ❼ return wrapper ➋ ➋ @my_decorator ⌂ def say_hello(): ⌃ print("Hello!") ⌄ ⌄ say_hello() ⌅ # Output ⌅ # Before function call ⌅ # Hello! ⌅ # After function call ⌆</pre>	
	<p>What built-in types are available in Python?</p> <p>Python provides several built-in data types. The main ones include:</p> <ol style="list-style-type: none"> 1. Numeric types: int, float, complex 2. Sequence types: list, tuple, range 3. Text type: str 4. Set types: set, frozenset 5. Mapping type: dict 6. Boolean type: bool 7. Binary types: bytes, bytearray, memoryview <p>These types allow you to store and manipulate different kinds of data.</p>		

	<p>How do you differentiate between .py and .pyc files in Python?</p> <p>"A .py file is the Python source code you write. A .pyc file is the compiled bytecode Python makes to run the program faster."</p>	<pre>python # my_module.py -> mã nguồn của bạn import my_module # Python sẽ tự động tạo my_module.pyc</pre>	<input type="button" value="Copy code"/>	
	<p>How do you create a Python function?</p> <p>"Use def to make a function, add a name and inputs, write the code inside, and use return to give back a value."</p>	<pre>python def add(a, b): return a + b result = add(3, 5) print(result) # Kết quả: 8</pre>	<input type="button" value="Copy code"/>	
	<p>How does a function return values in Python?</p> <p>"A Python function gives back a value using return. Without return, it gives back None."</p>			
	<p>What commands are used to delete Python files?</p> <p>In Python, you can delete files using the os or pathlib modules.</p> <p>Using os.remove():</p> <p>Using pathlib.Path.unlink():</p>	<pre>main.py > ... 1 # . Using os.remove() : 2 import os 3 os.remove("example.txt") # deletes the file 4 5 6 # . Using pathlib.Path.unlink() : 7 from pathlib import Path 8 file = Path("example.txt") 9 file.unlink() # deletes the file</pre>		
	<p>What are modules in Python?</p> <p>"A module is a Python file with code (functions, classes, variables) that you can import and use in other files."</p>	<pre>python # my_module.py def greet(name): print(f"Hello, {name}!") # main.py import my_module my_module.greet("Alice") # Output: Hello, Alice!</pre>	<input type="button" value="Copy code"/>	
	<p>What is a Python PATH?</p> <p>"Python PATH is a list of folders where Python looks for modules when you use import. It's in sys.path and you can add folders if needed."</p>	<pre>python import sys print(sys.path) # shows all directories Python searches for modules</pre>	<input type="button" value="Copy code"/>	

	<p>What is a Package in Python?</p> <p>"A package is a folder with Python modules and an <code>__init__.py</code> file that organizes them together."</p>	<p>markdown</p> <pre>my_package/ __init__.py module1.py module2.py</pre> <p>Bạn có thể import module từ package:</p> <p>python</p> <pre>from my_package import module1 module1.some_function()</pre>	
	<p>Create a module in Python.</p> <p>In Python, a module is simply a Python file with a <code>.py</code> extension that contains code such as functions, variables, or classes. You can then import it in another Python file to use its contents.</p> <p>"A module is a Python file with functions or classes that can be imported in another file."</p>	<p>1. Tạo file <code>my_module.py</code> với nội dung:</p> <p>python</p> <pre># my_module.py def greet(name): print(f"Xin chào, {name}!")</pre> <p>2. Sử dụng trong file khác:</p> <p>python</p> <pre># main.py import my_module my_module.greet("Alice") # Kết quả: Xin chào, Alice!</pre>	

	<p>What is lambda in Python?</p> <p>"A lambda is a small, anonymous Python function with only one expression that can take multiple arguments."</p>	<pre>python # Lambda function to add two numbers add = lambda x, y: x + y print(add(3, 5)) # Output: 8 # Using lambda with map numbers = [1, 2, 3] squared = list(map(lambda x: x**2, numbers)) print(squared) # Output: [1, 4, 9]</pre>	Copy code	
	<p>How memory can be managed in Python?</p> <p>"Python manages memory automatically using reference counting (tracks object use), garbage collection (removes unused objects), and dynamic allocation from its private heap. You usually don't need to free memory yourself."</p>			
	<p>What are keywords in Python?</p> <p>"Keywords are reserved words in Python that have special meanings and cannot be used as variable or function names."</p> <p>Keywords: if, else, for, while, def, class, import, return, try, except, finally,...</p>			
	<p>Is Python a case-sensitive language?</p> <p>"Yes, Python is case-sensitive, so names like myVar and myvar are different."</p>	<pre>python myVar = 10 myvar = 20 print(myVar) # Kết quả: 10 print(myvar) # Kết quả: 20</pre>	Copy code	
	<p>What are literals in Python?</p> <p>"A literal is a fixed value written in the code, like numbers (10), strings ("Hello"), or Booleans (True)."</p>	<pre>python x = 10 # numeric literal y = "Hello" # string literal flag = True # boolean literal</pre>	Copy code	
	<p>What is Type Conversion in Python?</p> <p>Type conversion is changing a value from one type to another using functions like int(), float(), or str().</p>	<pre>python x = 10 # int y = float(x) # convert int to float print(y) # Output: 10.0 z = str(x) # convert int to string print(z) # Output: "10"</pre>	Copy code	

	<p>What do you think is the use of dir () function in Python?</p> <p>"The dir() function shows all attributes and methods of an object, class, or module."</p>	<pre>python</pre> <div style="border: 1px solid #ccc; padding: 5px; display: inline-block;"> x = [1, 2, 3] print(dir(x)) # Lists all methods and attributes of a list object </div>	
	<p>How can you remove values from an array in Python?</p> <p>"In Python arrays, use remove() to delete a value or pop() to delete by index (last element if no index)."</p>	<pre>python</pre> <div style="border: 1px solid #ccc; padding: 5px; display: inline-block;"> import array arr = array.array('i', [1, 2, 3, 2]) arr.remove(2) # Xóa giá trị 2 đầu tiên print(arr) # Kết quả: array('i', [1, 3, 2]) arr.pop(1) # Xóa phần tử ở chỉ số 1 print(arr) # Kết quả: array('i', [1, 2]) </div>	
	How is list stored in memory in Python?		
	How is tuple stored in memory in Python?		
	<p>How are list, tuple, set, array, and dictionary stored in memory in Python?</p> <p>List, tuple, set, array và dictionary trong Python được lưu trữ trong bộ nhớ như thế nào?</p>		
Intermediate	<p>Is it possible for a function not to have a return statement and is it valid?</p> <p>"Python function can have no return. It is valid and will return None by default."</p>	<pre>python</pre> <div style="border: 1px solid #ccc; padding: 5px; display: inline-block;"> def greet(name): print(f"Hello, {name}!") result = greet("Alice") print(result) # Output: None </div>	

	<p>When should Python use triple quotes as a delimiter?</p> <p>"Use triple quotes in Python for multi-line strings or for docstrings to document modules, classes, or functions."</p>	<pre>python # Multi-line string text = """This is a string that spans multiple lines.""" # Docstring def greet(name): """This function greets the user by name.""" print(f"Hello, {name}!")</pre>	Copy code	
	<p>What is the main role of the <code>init</code> method? Give a code block example.</p> <p>The <code>__init__</code> method is used to set up a class's attributes when a new object is created.</p>	<pre>python class Person: def __init__(self, name, age): self.name = name self.age = age p = Person("Alice", 25) print(p.name) # Output: Alice print(p.age) # Output: 25</pre>	Copy code	
	<p>How do you convert string to lowercase in Python?</p> <p>In Python, you can convert a string to lowercase using the <code>lower()</code> method of the string.</p>	<pre>python text = "Hello, World!" lower_text = text.lower() print(lower_text) # Output: hello, world!</pre>	Copy code	

	<p>How do you use the split method in Python?</p> <p>In Python, the split() method is used to divide a string into a list of substrings based on a separator. By default, it splits by whitespace.</p>	<pre>python text = "Python is fun" words = text.split() # Split by whitespace print(words) # Output: ['Python', 'is', 'fun'] csv = "a,b,c" items = csv.split(",") # Split by comma print(items) # Output: ['a', 'b', 'c']</pre>	Copy code
	<p>What is a Try Block?</p> <p>A try block is used to run code that may cause an error. If an error happens, it can be handled in an except block so the program doesn't crash.</p>	<pre>main.py > ... 1 try: 2 x = 10 / 0 3 except ZeroDivisionError: 4 print("Cannot divide by zero!") 5 6 # Output 7 # Cannot divide by zero!</pre>	
	<p>What are generators in Python?</p> <p>A generator is a function that gives values one by one using yield, which helps save memory.</p>	<pre>main.py > ⓘ count_up_to 1 def count_up_to(n): 2 i = 1 3 while i <= n: 4 yield i 5 i += 1 6 7 for num in count_up_to(3): 8 print(num) 9 # Output: 10 # 1 11 # 2 12 # 3 13</pre>	

	<p>How can a module written in Python be accessed from C?</p> <p>"You can access a Python module from C by embedding the Python interpreter and using the Python/C API."</p>	<pre>main.cpp 1 #include <Python.h> 2 3 int main() { 4 // Start the Python interpreter 5 Py_Initialize(); 6 7 // Import Python module 8 PyObject *module = PyImport_ImportModule("mymodule"); 9 10 // Run Python code 11 PyRun_SimpleString("print('Accessed from C!')"); 12 13 // End the interpreter 14 Py_Finalize(); 15 return 0; 16} 17</pre>	
	<p>How a list can be reversed in Python?</p> <p>You can reverse a list by using <code>list.reverse()</code> (changes the list) or <code>list[::-1]</code> (makes a reversed copy).</p>	<p>python</p> <div style="text-align: right;">Copy code</div> <pre># Method 1: Using reverse() numbers = [1, 2, 3, 4] numbers.reverse() print(numbers) # Output: [4, 3, 2, 1] # Method 2: Using slicing nums = [1, 2, 3, 4] reversed_nums = nums[::-1] print(reversed_nums) # Output: [4, 3, 2, 1]</pre>	

	<p>What are ways to combine dataframes in Python?</p> <p>In Python (using pandas), there are several ways to combine DataFrames depending on how you want to merge the data:</p> <ol style="list-style-type: none"> 1. concat() – Combines DataFrames along rows or columns. 2. merge() – Combines DataFrames based on common columns or keys (like SQL joins). 3. join() – Combines DataFrames based on their indexes. <p>"We can combine DataFrames using concat() for stacking, merge() for SQL-style joins, and join() for index-based merging."</p>	<pre><code>python import pandas as pd df1 = pd.DataFrame({'A': [1, 2], 'B': [3, 4]}) df2 = pd.DataFrame({'A': [5, 6], 'B': [7, 8]}) # 1. concat() combined = pd.concat([df1, df2]) print(combined) # 2. merge() df3 = pd.DataFrame({'A': [1, 2], 'C': ['x', 'y']}) merged = pd.merge(df1, df3, on='A') print(merged) # 3. join() joined = df1.join(df3.set_index('A'), on='A') print(joined)</code></pre>	 Copy code
	<p>What is a PIP?</p> <p>"PIP is Python's package installer, used to install and manage third-party libraries."</p>	<pre><code>bash # Cài đặt package pip install requests # Nâng cấp package pip install --upgrade requests # Liệt kê các package đã cài pip list</code></pre>	

	<p>Why is finalize used in Python?</p> <p>In Python, finalize (from the weakref module) is used to register a callback function that will be called when an object is about to be garbage collected. It is helpful for cleaning up resources like closing files or network connections when an object is no longer needed, without relying on <code>__del__()</code>.</p> <p>"finalize is used to register a callback for cleanup when an object is about to be garbage collected."</p>	<pre>main.py > ... 1 import weakref 2 3 class MyClass: 4 def __init__(self, name): 5 self.name = name 6 7 def goodbye(obj_name): 8 print(f"{obj_name} is being finalized") 9 10 obj = MyClass("TestObject") 11 finalizer = weakref.finalize(obj, goodbye, obj.name) 12 13 del obj # Triggers the finalize callback 14 15 # Output: 16 # TestObject is being finalized 17</pre>
	<p>Differentiate between override and new modifiers.</p> <p>The difference between override and new modifiers (common in languages like C#; Python doesn't have explicit modifiers but conceptually similar behavior can exist) is:</p> <ol style="list-style-type: none"> 1. Override: <ul style="list-style-type: none"> - Used when a derived class wants to provide a new implementation for a method already defined in the base class. - Ensures polymorphic behavior, so the derived class method is called even when using a base class reference. 2. New: <ul style="list-style-type: none"> - Used when a derived class defines a method with the same name as in the base class but does not override it. - It hides the base class method; the base class version is called when the object is referenced as the base type. 	
	<p>In Python, what would you do to create an empty class?</p> <p>In Python, to create an empty class, you define the class with the <code>class</code> keyword and use <code>pass</code> inside to indicate no content.</p>	<pre>python class EmptyClass: pass obj = EmptyClass() print(obj) # Output: <__main__.EmptyClass object at 0x...></pre> <div style="text-align: right;"> Copy code </div>

	<p>Do you think you can call the parent class without its instance creation?</p> <p>In Python, you can call methods or access attributes of a parent class without creating an instance of it if the method is a class method (@classmethod) or a static method (@staticmethod). For regular instance methods, you normally need an instance of the parent class.</p> <p>"You can call a parent class without an instance only if the method is a class method or static method. For normal methods, you need an object."</p>	<pre><code>python</code></pre> <pre><code>class Parent: @staticmethod def greet(): print("Hello from Parent") # Call static method without creating instance Parent.greet() # Output: Hello from Parent</code></pre>	Copy code
	<p>In what ways can parent members in a child class be accessed?</p> <p>In Python, parent class members (attributes or methods) can be accessed in a child class in several ways:</p> <ol style="list-style-type: none"> 1. Using super() – Calls parent methods or accesses parent attributes in a safe and flexible way. 2. Directly using the parent class name – Calls parent methods or accesses attributes explicitly. 3. Inherited automatically – If the member is public, the child class can use it directly. 	<pre><code>main.py > ... 1 class Parent: 2 def __init__(self): 3 self.value = 10 4 5 def show(self): 6 print("Parent value:", self.value) 7 8 class Child(Parent): 9 def display(self): 10 # Using super() 11 super().show() 12 # Using parent class name directly 13 Parent.show(self) 14 # Direct access to inherited attribute 15 print("Access value directly:", self.value) 16 17 c = Child() 18 c.display() 19 20 # Parent value: 10 21 # Parent value: 10 22 # Access value directly: 10</code></pre>	

	<p>What are Pandas in Python?</p> <p>Pandas is a Python library used for data manipulation and analysis. It provides data structures like DataFrame and Series that make it easy to handle tabular or time-series data, perform filtering, aggregation, and statistical operations efficiently.</p> <p>"Pandas is a Python library for data manipulation and analysis, using DataFrame and Series structures."</p>	<pre>python import pandas as pd data = {'Name': ['Alice', 'Bob'], 'Age': [25, 30]} df = pd.DataFrame(data) print(df) markdown</pre> <table border="1"> <thead> <tr> <th></th> <th>Name</th> <th>Age</th> </tr> </thead> <tbody> <tr> <td>0</td> <td>Alice</td> <td>25</td> </tr> <tr> <td>1</td> <td>Bob</td> <td>30</td> </tr> </tbody> </table>		Name	Age	0	Alice	25	1	Bob	30	
	Name	Age										
0	Alice	25										
1	Bob	30										
	<p>What is a NumPy?</p> <p>NumPy (Numerical Python) is a Python library for numerical computing. It provides a high-performance multidimensional array object (ndarray) and functions to perform mathematical, logical, and statistical operations efficiently. It is widely used in scientific computing, data analysis, and machine learning.</p> <p>"NumPy is a Python library for numerical computing with high-performance arrays and mathematical functions."</p>	<pre>python import numpy as np arr = np.array([1, 2, 3, 4]) print(arr * 2) # Output: [2 4 6 8]</pre>										
	<p>Why is NumPy preferred over Python lists?</p> <p>NumPy is preferred over Python lists for several reasons:</p> <ol style="list-style-type: none"> 1. Performance: NumPy arrays are implemented in C and are much faster than Python lists for large-scale numerical computations. 2. Memory efficiency: NumPy arrays use less memory because they store elements in a homogeneous and contiguous block. 3. Vectorized operations: NumPy allows element-wise operations without explicit loops, making code cleaner and faster. 4. Rich mathematical functions: NumPy provides a wide range of built-in mathematical, statistical, and linear algebra functions. <p>"NumPy is faster, more memory-efficient, supports vectorized operations, and has rich mathematical functions compared to Python lists."</p>	<pre>python import numpy as np arr = np.array([1, 2, 3, 4]) print(arr * 2) # Output: [2 4 6 8] (vectorized operation)</pre>										

	<p>How can we efficiently load data from a text file?</p> <p>In Python, data from a text file can be efficiently loaded depending on the file size and format:</p> <ol style="list-style-type: none"> 1. Using <code>open()</code> – Reads the file safely and automatically closes it. 2. Using <code>readlines()</code> or iterating line by line – Efficient for large files to avoid loading everything into memory. 3. Using <code>pandas.read_csv()</code> – For structured data (like CSV or tab-delimited), this is fast and convenient. 4. Using <code>numpy.loadtxt()</code> – For numerical data, it's efficient and returns NumPy arrays. 	<pre>python # Method 1: Using open() and iteration with open("data.txt") as f: for line in f: print(line.strip()) # Method 2: Using pandas import pandas as pd df = pd.read_csv("data.txt", delimiter="\t") print(df # Method 3: Using NumPy import numpy as np data = np.loadtxt("data.txt") print(data)</pre>	 Copy code
	<p>What is reindexing in Pandas?</p> <p>In Pandas, reindexing means changing the row or column labels of a DataFrame or Series to a new set of labels. This can be used to add new rows/columns, remove some, or rearrange the order. When new labels are introduced, the missing values are filled with NaN by default or a specified value.</p> <p>"Reindexing in Pandas changes the row or column labels, adding or rearranging them and filling missing values."</p>	<pre>python import pandas as pd df = pd.DataFrame({'A': [1, 2, 3]}, index=[0, 1, 2]) new_index = [0, 1, 2, 3] # Adding a new row df_reindexed = df.reindex(new_index, fill_value=0) print(df_reindexed) css A 0 1 1 2 2 3 3 0</pre>	 Copy code

	<p>How can you copy an object in Python?</p> <p>In Python, you can copy an object in two main ways:</p> <ol style="list-style-type: none"> 1. Shallow copy – Creates a new object, but nested objects are shared. Use <code>copy.copy()</code>. 2. Deep copy – Creates a new object with all nested objects copied recursively. Use <code>copy.deepcopy()</code>. <p>"Use <code>copy.copy()</code> for shallow copy and <code>copy.deepcopy()</code> for deep copy of objects."</p>	<pre>python import copy original = [[1, 2], [3, 4]]</pre> <p># Shallow copy</p> <pre>shallow = copy.copy(original) shallow[0][0] = 100 print(original) # Output: [[100, 2], [3, 4]]</pre> <p># Deep copy</p> <pre>deep = copy.deepcopy(original) deep[0][0] = 999 print(original) # Output: [[100, 2], [3, 4]] (unchanged)</pre>	Copy code
	<p>What is shallow and deep copying in Python?</p> <p>The difference between shallow copy and deep copy in Python is:</p> <ol style="list-style-type: none"> 1. Shallow copy: <ul style="list-style-type: none"> - Creates a new object, but the nested objects inside are shared with the original. - Changes to nested objects affect both the original and the copied object. - Use <code>copy.copy()</code>. 2. Deep copy: <ul style="list-style-type: none"> - Creates a completely independent copy, including all nested objects recursively. - Changes in the copied object do not affect the original object. - Use <code>copy.deepcopy()</code>. <p>"Shallow copy copies only the top-level object, sharing nested objects, while deep copy copies everything recursively, making an independent object."</p>	<pre>python import copy original = [[1, 2], [3, 4]]</pre> <p>shallow = copy.copy(original) shallow[0][0] = 100 print(original) # Output: [[100, 2], [3, 4]] (nested changed)</p> <p>deep = copy.deepcopy(original) deep[0][0] = 999 print(original) # Output: [[100, 2], [3, 4]] (original unchanged)</p>	Copy code

	<p>What are Pickling?</p> <p>In Python, pickling is the process of converting a Python object into a byte stream so that it can be saved to a file or sent over a network. The reverse process, converting the byte stream back to a Python object, is called unpickling. Pickling is commonly used for persisting Python objects.</p> <p>"Pickling is serializing a Python object to a byte stream for storage or transmission; unpickling restores it."</p>	<pre>main.py > ... 1 import pickle 2 3 data = {'name': 'Alice', 'age': 25} 4 5 # Pickling (serialize) 6 with open('data.pkl', 'wb') as f: 7 pickle.dump(data, f) 8 9 # Unpickling (deserialize) 10 with open('data.pkl', 'rb') as f: 11 loaded_data = pickle.load(f) 12 13 print(loaded_data) 14 15 # output 16 # {'name': "Alice", 'age' : 25} 17</pre>	
	<p>How can you define Unpickling in Python?</p> <p>In Python, unpickling is the process of converting a pickled (serialized) byte stream back into the original Python object. It is the reverse of pickling and is used to load objects that were previously saved or transmitted.</p> <p>"Unpickling converts a pickled byte stream back into the original Python object."</p>	<p>python</p> <pre>import pickle # Assume 'data.pkl' was created by pickling with open('data.pkl', 'rb') as f: loaded_data = pickle.load(f) print(loaded_data)</pre> <p>bash</p> <pre>{'name': 'Alice', 'age': 25}</pre>	Copy code Copy code

	<p>What function is used for Pickling and Unpickling?</p> <p>In Python, the pickle module provides functions for pickling and unpickling:</p> <ol style="list-style-type: none"> 1. Pickling (serialization): Use pickle.dump() to write an object to a file or pickle.dumps() to get a byte stream. 2. Unpickling (deserialization): Use pickle.load() to read an object from a file or pickle.loads() to load from a byte stream. 	<pre>main.py > ... 1 import pickle 2 3 data = {'name': 'Alice', 'age' : 25} 4 5 # Pickling to a file 6 with open('data.pkl', 'wb') as f: 7 pickle.dump(data, f) 8 9 # Unpickling from a file 10 with open('data.pkl', 'rb') as f: 11 loaded_data = pickle.load(f) 12 13 print(loaded_data) 14 15 # output 16 # {'name': 'Alice', 'age': 25} 17</pre>	
	<p>What are some types of Type Conversion in Python?</p> <p>In Python, type conversion is changing a variable from one data type to another. There are two main types:</p> <ol style="list-style-type: none"> 1. Implicit Type Conversion (Type Casting): Python automatically converts one type to another. Example: converting int to float. 2. Explicit Type Conversion: You manually convert types using built-in functions like: <ul style="list-style-type: none"> - int() – converts to integer - float() – converts to float - str() – converts to string - list() – converts to list - tuple() – converts to tuple - dict() – converts to dictionary 	<pre>python Copy code x = 10 # int y = float(x) # explicit conversion to float print(y) # Output: 10.0</pre>	
	<p>Give an example of a lambda function.</p> <p>A lambda function in Python is an anonymous, single-line function defined using the lambda keyword.</p> <p>"A lambda function is an anonymous single-line function defined with lambda."</p>	<pre>python Copy code # Lambda function to add two numbers add = lambda x, y: x + y print(add(5, 3)) # Output: 8</pre>	

	<p>In Python, what is Polymorphism?</p> <p>In Python, polymorphism is the ability of different objects to be accessed through the same interface, allowing them to perform actions in different ways depending on their type. It is commonly seen in method overriding, operator overloading, and duck typing.</p> <p>"Polymorphism allows different objects to be accessed through the same interface, performing actions in their own way."</p>	<pre>main.py > ... 1 class Cat: 2 def speak(self): 3 print("Meow") 4 5 class Dog: 6 def speak(self): 7 print("Woof") 8 9 def make_animal_speak(animal): 10 animal. speak() 11 12 cat = Cat() 13 dog = Dog() 14 15 make_animal_speak(cat) # Output: Meow 16 make_animal_speak(dog) # Output: Woof 17</pre>	
	<p>How would you define the Swapcase () in Python?</p> <p>In Python, the swapcase() method is a string method that converts all uppercase letters to lowercase and all lowercase letters to uppercase in a string. It does not change non-alphabetic characters.</p> <p>"swapcase() converts uppercase letters to lowercase and lowercase letters to uppercase in a string."</p>	<pre>python text = "Hello World" swapped = text.swapcase() print(swapped) # Output: hELLO wORLD</pre>	Copy code
Advanced	<p>What does [::-1] do in Python and give an example?</p> <p>In Python, the [::-1] slice notation is used to reverse a sequence such as a string, list, or tuple. It means: start from the end towards the first element with a step of -1.</p> <p>"[::-1] reverses a sequence in Python, like a string or list."</p>	<pre>python text = "Hello" reversed_text = text[::-1] print(reversed_text) # Output: olleH numbers = [1, 2, 3, 4] print(numbers[::-1]) # Output: [4, 3, 2, 1]</pre>	Copy code

	<p>Explain database connection in Python Flask.</p> <p>In Flask, connecting to a database involves using a database driver or ORM (Object Relational Mapper). Common approaches include:</p> <ol style="list-style-type: none"> 1. Using Flask extensions like Flask-SQLAlchemy: <ul style="list-style-type: none"> - Provides ORM support and simplifies database operations. - You define the database URI, create a SQLAlchemy object, and interact with the database through models. 2. Using direct database drivers (like sqlite3 or pymysql): <ul style="list-style-type: none"> - You manually open a connection, create a cursor, execute queries, and close the connection. 	<pre>python from flask import Flask from flask_sqlalchemy import SQLAlchemy app = Flask(__name__) app.config['SQLALCHEMY_DATABASE_URI'] = 'sqlite:///test.db' db = SQLAlchemy(app) class User(db.Model): id = db.Column(db.Integer, primary_key=True) name = db.Column(db.String(50)) # Create tables with app.app_context(): db.create_all()</pre>	
	<p>What is the Dogpile effect?</p> <p>The Dogpile effect occurs when multiple processes or threads try to regenerate the same cached data simultaneously after it expires. This can overload the database or backend system because all requests trigger the same expensive computation at once.</p> <p>Prevention:</p> <ul style="list-style-type: none"> - Use cache locks or stale-while-revalidate techniques to ensure only one process regenerates the cache while others wait. <p>Example (conceptual):</p> <p>Imagine a web cache for a popular product price. When the cache expires, hundreds of users request the price at the same time, causing many queries to the database simultaneously — this is the Dogpile effect.</p> <p><i>"Dogpile effect happens when multiple requests try to regenerate expired cache simultaneously, causing overload."</i></p>		

	<p>Are multiple inheritances supported in Python?</p> <p>Yes, Python supports multiple inheritance, which means a class can inherit attributes and methods from more than one parent class. Python uses the Method Resolution Order (MRO) to determine the order in which base classes are searched when calling a method.</p> <p>"Yes, Python supports multiple inheritance and uses MRO to determine method lookup order."</p>	<pre>python class Parent1: def greet(self): print("Hello from Parent1") class Parent2: def greet(self): print("Hello from Parent2") class Child(Parent1, Parent2): pass c = Child() c.greet() # Output: Hello from Parent1 (MRO decides Parent1 first)</pre>	
	<p>Does Python make use of access specifiers?</p> <p>Python does not have strict access specifiers like public, private, or protected found in languages such as Java or C++. Instead, it relies on naming conventions:</p> <ol style="list-style-type: none"> 1. Public: Any attribute or method without underscores is public and accessible anywhere. 2. Protected: Prefix a single underscore _ (e.g., _var) to indicate it should be treated as protected, but it's only a convention. 3. Private: Prefix a double underscore __ (e.g., __var) triggers name mangling, making it harder (but not impossible) to access from outside the class. <p>"Python uses naming conventions instead of strict access specifiers: public, protected (_), and private (__ with name mangling)."</p>	<pre>main.py > ... 1 class MyClass: 2 public_var = 1 3 _protected_var = 2 4 __private_var = 3 5 6 obj = MyClass() 7 8 # Accessible 9 print(obj.public_var) 10 11 # Accessible, but should be treated as protected 12 print(obj._protected_var) 13 14 # Raises AttributeError 15 # print(obj.__private_var) 16 17 # Access private using name mangling 18 print(obj._MyClass__private_var) 19 </pre>	

	<p>How do you create a constructor in Python?</p> <p>In Python, a constructor is a special method called <code>__init__()</code> that is automatically invoked when an object of a class is created. It is used to initialize the attributes of the object.</p> <p>"A constructor in Python is the <code>__init__()</code> method used to initialize object attributes."</p>	<pre>python class Person: def __init__(self, name, age): self.name = name self.age = age p = Person("Alice", 25) print(p.name) # Output: Alice print(p.age) # Output: 25</pre>	<input type="button" value="Copy code"/>
	<p>How to save an image in Python locally when we know the URL address?</p> <p>In Python, you can save an image from a URL using the <code>requests</code> library or <code>urllib</code>. The idea is to download the image content and write it to a local file in binary mode.</p> <p>"Use <code>requests.get()</code> or <code>urllib.request.urlretrieve()</code> to download an image from a URL and save it locally."</p>	<pre>main.py > ... 1 import requests 2 3 url = "https://example.com/image.jpg" 4 response = requests.get(url) 5 6 if response.status_code == 200: 7 with open("local_image.jpg", "wb") as f: 8 f.write(response.content) 9 10 import urllib.request 11 12 url = "https://example.com/image.jpg" 13 urllib.request.urlretrieve(url, "local_image.jpg") 14</pre>	
	<p>Explain how the <code>join()</code> function in Python?</p> <p>The <code>join()</code> function in Python is a string method used to combine (concatenate) elements of an iterable (like a list or tuple) into a single string, with a specified separator between each element.</p>	<pre>python words = ['Python', 'is', 'fun'] sentence = ' '.join(words) print(sentence) kotlin Python is fun</pre>	<input type="button" value="Copy code"/>

	<p>How can you identify and deal with missing values in a dataframe?</p> <p>In Python's Pandas, missing values are usually represented as NaN (Not a Number). You can identify and handle them using built-in functions.</p>	<pre>❖ main.py > ... 1 import pandas as pd 2 3 df = pd.DataFrame({ 4 'Name': ['Alice', 'Bob', None], 5 'Age': [25, None, 30] 6 }) 7 8 # 1. Identify missing values: 9 print(df.isnull()) # Shows True where values are missing 10 print(df.isnull().sum()) # Counts total missing values per column 11 12 # 2. Handle missing values: 13 # Remove missing values: 14 df = df.dropna() 15 16 # Fill missing values: 17 df = df.fillna({'Name': 'Unknown', 'Age': df['Age'].mean()}) 18</pre>	
	<p>What is the use of manage.py in Python?</p> <p>In Django (a popular Python web framework), the manage.py file is an auto-generated command-line utility that helps you interact with your Django project.</p> <p>It allows you to perform various administrative tasks, such as:</p> <ul style="list-style-type: none"> - Running the development server (python manage.py runserver) - Applying database migrations (python manage.py migrate) - Creating new apps (python manage.py startapp appname) - Creating superusers (python manage.py createsuperuser) - Running tests (python manage.py test) <p>"manage.py acts as a wrapper around Django's django-admin tool and helps manage your project more easily."</p>		
	<p>Explain the shuffle method and give an example.</p> <p>The shuffle() method in Python is part of the random module. It is used to randomly rearrange (shuffle) the elements of a list in place — meaning it modifies the original list instead of creating a new one.</p>	<pre>python import random numbers = [1, 2, 3, 4, 5] random.shuffle(numbers) print(numbers)</pre>	Copy code

	<p>What method can be used to generate random numbers in Python?</p> <p>In Python, you can generate random numbers using functions from the random module. Some commonly used methods are:</p> <ol style="list-style-type: none"> 1. random.random() – Returns a random float between 0.0 and 1.0. 2. random.randint(a, b) – Returns a random integer between a and b (inclusive). 3. random.uniform(a, b) – Returns a random float between a and b. 4. random.choice(sequence) – Returns a random element from a sequence (like list or tuple). 	<pre>python import random print(random.random()) # Output: 0.374 (example) print(random.randint(1, 10)) # Output: 7 (example) print(random.uniform(1, 5)) # Output: 3.456 (example) print(random.choice([10, 20, 30])) # Output: 20 (example)</pre>	Copy code
	<p>What does *args, **kwargs mean in Python?</p> <p>In Python, *args and **kwargs are used in function definitions to allow variable numbers of arguments.</p> <ol style="list-style-type: none"> 1. *args (Non-keyword arguments): <ul style="list-style-type: none"> - Collects extra positional arguments passed to a function into a tuple. - Useful when you don't know in advance how many positional arguments will be provided. 2. **kwargs (Keyword arguments): <ul style="list-style-type: none"> - Collects extra keyword arguments passed to a function into a dictionary. - Useful when you want to handle named arguments dynamically. 	<pre>python def demo(*args, **kwargs): print("args:", args) print("kwargs:", kwargs) demo(1, 2, 3, name="Alice", age=25) yaml args: (1, 2, 3) kwargs: {'name': 'Alice', 'age': 25}</pre>	Copy code
	<p>Is Flask an MVC model? If true, justify this using the MVC pattern.</p> <p>Flask can be used to implement the MVC (Model-View-Controller) pattern, but it is not strictly enforced. Flask is a micro-framework, so it gives developers flexibility to structure their application in an MVC way.</p> <ul style="list-style-type: none"> - Model: Handles data and business logic. In Flask, this is usually done using ORMs like SQLAlchemy or by directly interacting with databases. - View: Represents the user interface, typically HTML templates rendered with Jinja2 in Flask. - Controller: Handles application logic and routing, implemented via Flask route functions that take requests, process data (using models), and return rendered views. 	<pre>graphql app/ ├── models.py # Model: database tables and business logic ├── templates/ # View: HTML templates (Jinja2) ├── routes.py # Controller: routes and request handling └── app.py # Main application</pre>	Copy code

	<p>How can we check if all characters in a string are alphanumeric?</p> <p>In Python, you can use the <code>isalnum()</code> string method to check if all characters in a string are alphanumeric (letters and numbers only, no spaces or special characters). It returns True if all characters are alphanumeric, otherwise False.</p> <p>"Use the <code>isalnum()</code> method; it returns True if all characters in a string are letters or digits."</p>	<pre>python text1 = "Python123" text2 = "Python 123" print(text1.isalnum()) # Output: True print(text2.isalnum()) # Output: False (space is not alphanumeric)</pre>	
	<p>What are some of the most used built-in modules in Python?</p> <p>Python comes with many built-in modules that provide useful functionality without requiring installation. Some of the most commonly used built-in modules are:</p> <ol style="list-style-type: none"> 1. <code>os</code> – Interact with the operating system (files, directories, environment variables). 2. <code>sys</code> – Access system-specific parameters and functions. 3. <code>math</code> – Perform mathematical operations (trigonometry, logarithms, etc.). 4. <code>random</code> – Generate random numbers and perform random operations. 5. <code>datetime</code> – Work with dates and times. 6. <code>json</code> – Encode and decode JSON data. 7. <code>re</code> – Perform regular expression operations. 8. <code>collections</code> – Specialized container datatypes like <code>Counter</code>, <code>deque</code>, <code>OrderedDict</code>. <p>"Common built-in modules include <code>os</code>, <code>sys</code>, <code>math</code>, <code>random</code>, <code>datetime</code>, <code>json</code>, <code>re</code>, and <code>collections</code>."</p>		

	<p>What are the tools for debugging and performing static analysis in Python?</p> <p>In Python, there are several tools for debugging and static analysis:</p> <ol style="list-style-type: none"> 1. Debugging tools: <ul style="list-style-type: none"> - pdb – The built-in Python debugger, allows stepping through code, setting breakpoints, and inspecting variables. - ipdb – An enhanced version of pdb with IPython integration. - IDE debuggers – Most IDEs like PyCharm, VS Code, and Spyder provide integrated debugging with breakpoints, watches, and call stack inspection. 2. Static analysis tools: <ul style="list-style-type: none"> - pylint – Checks for coding standards, errors, and potential bugs. - flake8 – Linter for style guide enforcement (PEP 8) and code errors. - mypy – Performs type checking based on type hints. - pyflakes – Detects errors without performing style checks. <p><i>"For debugging use pdb, ipdb, or IDE debuggers; for static analysis use pylint, flake8, mypy, or pyflakes."</i></p>	
	<p>When will the else part of try-except-else be executed?</p> <p>"The else part runs only when the try block finishes with no errors."</p>	<pre>python try: result = 10 / 2 except ZeroDivisionError: print("Cannot divide by zero!") else: print("Division succeeded, result is", result) csharp Division succeeded, result is 5.0</pre>

	<p>Write a code to swap two numbers in Python.</p> <p>In Python, you can swap two numbers in a very simple way using tuple unpacking, without needing a temporary variable.</p> <p>"You can swap two numbers using <code>a, b = b, a</code> or with a temporary variable."</p>	<pre>❖ main.py > ... 1 2 a = 5 3 b = 10 4 5 # Tuple unpacking 6 a, b = b, a # Swap 7 print("a =", a) # Output: a = 10 8 print("b =", b) # Output: b = 5 9 10 # Temporary variable 11 temp = a 12 a = b 13 b = temp 14</pre>	
	<p>Show code examples of deep and shallow copying.</p> <p>In Python, shallow copy and deep copy are used to copy objects, but they behave differently:</p> <ul style="list-style-type: none"> - Shallow copy: Copies the object itself, but nested objects are still references to the original. - Deep copy: Copies the object and all nested objects, creating a completely independent copy. <p>"Shallow copy copies the outer object but keeps references to nested objects, while deep copy duplicates everything including nested objects."</p>	<pre>❖ main.py > ... 1 import copy 2 3 # Danh sách gốc với danh sách lồng 4 original = [1, 2, [3, 4]] 5 6 # Shallow copy 7 shallow = copy.copy(original) 8 shallow[2][0] = 99 9 print("Original sau shallow copy:", original) # Danh sách lồng bị thay đổi 10 # Output: 11 # Original sau shallow copy: [1, 2, [99, 4]] 12 13 # Deep copy 14 deep = copy.deepcopy(original) 15 deep[2][1] = 88 16 print("Original sau deep copy:", original) # Danh sách gốc không thay đổi 17 # Output: 18 # Original sau deep copy: [1, 2, [99, 4]]</pre>	
	<p>Write a code to reverse multiple values from functions.</p> <p>In Python, a function can return multiple values as a tuple, and you can reverse them using tuple unpacking or the <code>[::-1]</code> slicing.</p> <p>"Return multiple values as a tuple and use slicing <code>[::-1]</code> to reverse them."</p>	<p>python</p> <p>Copy code</p> <pre>def get_values(): return 1, 2, 3, 4 # Get values and reverse them a, b, c, d = get_values()[::-1] print(a, b, c, d) # Output: 4 3 2 1</pre>	

	<p>What are the steps to create 3D, 2D, and 1D arrays?</p> <p>In Python, you can create 1D, 2D, and 3D arrays using the NumPy library.</p>	<pre>➊ main.py > ... 1 import numpy as np 2 3 array_1d = np.array([1, 2, 3, 4]) 4 print("Mảng 1D:", array_1d) 5 6 array_2d = np.array([[1, 2], [3, 4]]) 7 print("Mảng 2D:\n", array_2d) 8 9 array_3d = np.array([[[1, 2], [3, 4]], [[5, 6], [7, 8]]]) 10 print("Mảng 3D:\n", array_3d) 11</pre>	
	<p>Write a code to check whether two words are anagrams or not.</p> <p>In Python, two words are anagrams if they contain the same letters with the same frequency, regardless of order. One simple way is to sort the letters and compare.</p>	<pre>➊ main.py > ... 1 def are_anagrams(word1, word2): 2 return sorted(word1.lower()) == sorted(word2.lower()) 3 4 # Test 5 print(are_anagrams("listen", "silent")) # Output: True 6 print(are_anagrams("hello", "world")) # Output: False 7 8 9 # Another way is to use collections.Counter 10 from collections import Counter 11 12 def are_anagrams(word1, word2): 13 return Counter(word1.lower()) == Counter(word2.lower()) 14</pre>	
	<p>Write a code to test for the mode in a list of numbers.</p> <p>In Python, the mode of a list is the number that appears most frequently. You can use the statistics module to find it.</p> <p>Note: If multiple numbers have the same highest frequency, statistics.mode() will return the first one encountered.</p>	<pre>➊ main.py > ... 1 from statistics import mode 2 3 numbers = [1, 2, 2, 3, 4, 2, 5, 3, 3] 4 5 # Tìm mode 6 most_frequent = mode(numbers) 7 print("Mode của danh sách là:", most_frequent) 8 9 # Output: 10 # Mode của danh sách là: 3 11</pre>	

	<p>Show how to access the dataset of a publicly shared spreadsheet in the format of CSV stored in Google Drive.</p> <p>To access a publicly shared Google Spreadsheet in CSV format, you can use Python's pandas library with the file's CSV export link.</p> <p>Steps:</p> <ol style="list-style-type: none"> 1. Make sure the spreadsheet is public. 2. Get the CSV export link. For a Google Sheet with URL like: https://docs.google.com/spreadsheets/d/FILE_ID/edit#gid=0 The CSV link will be: https://docs.google.com/spreadsheets/d/FILE_ID/export?format=csv 3. Use <code>pandas.read_csv()</code> to load it: 	<pre>python import pandas as pd url = "https://docs.google.com/spreadsheets/d/FILE_ID/export?format=csv" df = pd.read_csv(url) print(df.head())</pre> <p>This will load the spreadsheet as a DataFrame, allowing you to analyze it in Python.</p>
	<p>Write a code to add two integers without using the '+' operator when the numbers are greater than zero.</p> <p>You can add two positive integers without using the + operator by using bitwise operations. The idea is to use XOR (^) for addition without carry and AND (&) shifted left for the carry, repeating until there is no carry left.</p> <p>"Use bitwise XOR for addition and AND shifted left for carry, loop until carry is zero."</p>	<pre>python def add_without_plus(a, b): while b != 0: carry = a & b # Find carry a = a ^ b # Add without carry b = carry << 1 # Shift carry return a # Test x = 7 y = 5 print(add_without_plus(x, y)) # Output: 12</pre>

	<p>Write a code that is given a sequence of numbers and it checks if the numbers are unique.</p> <p>In Python, you can check if numbers in a sequence are unique by converting the sequence to a set and comparing its length with the original sequence.</p> <p>"Convert the sequence to a set and compare lengths, or check each number with a set to ensure uniqueness."</p>	<pre>⌚ main.py > ... 1 def are_numbers_unique(numbers): 2 return len(numbers) == len(set(numbers)) 3 4 # Kiểm tra 5 nums1 = [1, 2, 3, 4, 5] 6 nums2 = [1, 2, 2, 3, 4] 7 8 print(are_numbers_unique(nums1)) # Kết quả: True 9 print(are_numbers_unique(nums2)) # Kết quả: False 10 11 # Cách khác dùng vòng lặp: 12 def are_numbers_unique(numbers): 13 seen = set() 14 for num in numbers: 15 if num in seen: 16 return False 17 seen.add(num) 18 return True 19</pre>	
	<p>Write a program to convert the date format from yyyy-mm-dd to dd-mm-yyyy.</p> <p>In Python, you can convert a date string from yyyy-mm-dd to dd-mm-yyyy using the `datetime` module.</p> <p>"Use `datetime.strptime()` to parse yyyy-mm-dd, then `strftime()` to format dd-mm-yyyy."</p>	<p>python</p> <div style="float: right;">Copy code</div> <pre>from datetime import datetime date_str = "2025-11-13" # Original format yyyy-mm-dd # Convert to datetime object date_obj = datetime.strptime(date_str, "%Y-%m-%d") # Format to dd-mm-yyyy new_date_str = date_obj.strftime("%d-%m-%Y") print(new_date_str) # Output: 13-11-2025</pre>	

	<p>Write a program in Python to create a Fibonacci series.</p> <p>In Python, you can create a Fibonacci series using a loop or recursion.</p> <p>"Start with [0,1] and use a loop to append the sum of the last two numbers to build the Fibonacci series."</p>	<pre><code>main.py > ... 1 def fibonacci(n): 2 fib_series = [0, 1] # First two numbers 3 for i in range(2, n): 4 fib_series.append(fib_series[i-1] + fib_series[i-2]) 5 return fib_series[:n] 6 7 # Test 8 n = 10 9 print(fibonacci(n)) 10 11 # output 12 # [0, 1, 1, 2, 3, 5, 8, 13, 21, 34]</code></pre>	
	<p>Write a code to create a single string from elements in a list.</p> <p>In Python, you can create a single string from a list of elements using the join() method.</p>	<p>python</p> <p><code>words = ["Python", "is", "fun"]</code></p> <p><code># Combine list elements into a single string with spaces</code></p> <p><code>single_string = " ".join(words)</code></p> <p><code>print(single_string) # Output: Python is fun</code></p> <p><code># Combine without spaces</code></p> <p><code>single_string_no_space = "".join(words)</code></p> <p><code>print(single_string_no_space) # Output: Pythonisfun</code></p>	Copy code

	<p>Write a program to check if a number is a prime.</p> <p>In Python, a prime number is a number greater than 1 that has no divisors other than 1 and itself. You can check it using a simple loop.</p>	<pre>python def is_prime(n): if n <= 1: return False for i in range(2, int(n**0.5) + 1): if n % i == 0: return False return True # Test num = 29 print(f"{num} is prime? {is_prime(num)}") # Output: True</pre>	Copy code
	<p>Write a program to calculate the median in Python using the NumPy arrays.</p> <p>In Python, you can calculate the median of a dataset using NumPy with the np.median() function.</p>	<pre>python import numpy as np # Sample NumPy array data = np.array([10, 20, 30, 40, 50]) # Calculate median median_value = np.median(data) print("Median is:", median_value) # Output: 30.0</pre>	Copy code

	<p>Write a program to execute the interchange sort algorithm.</p> <p>Interchange Sort is a simple sorting algorithm similar to bubble sort, where every element is compared with all other elements and swapped if they are out of order.</p> <p>"Compare each element with all subsequent elements and swap if out of order."</p>	<pre>main.py > ... 1 def interchange_sort(arr): 2 n = len(arr) 3 for i in range(n-1): 4 for j in range(i+1, n): 5 if arr[i] > arr[j]: 6 arr[i], arr[j] = arr[j], arr[i] 7 return arr 8 9 # Test 10 numbers = [29, 10, 14, 37, 13] 11 sorted_numbers = interchange_sort(numbers) 12 print("Sorted array:", sorted_numbers) 13 14 # Output: 15 # Sorted array: [10, 13, 14, 29, 37] 16</pre>	
	<p>Write a program to execute the selection sort algorithm.</p> <p>Selection Sort is a sorting algorithm where the smallest (or largest) element is repeatedly selected from the unsorted portion and swapped with the first unsorted element.</p> <p>"Repeatedly select the minimum element from the unsorted part and swap it with the first unsorted element."</p>	<pre>main.py > ... 1 def selection_sort(arr): 2 n = len(arr) 3 for i in range(n): 4 # Assume the minimum is the first unsorted element 5 min_idx = i 6 for j in range(i+1, n): 7 if arr[j] < arr[min_idx]: 8 min_idx = j 9 # Swap the found minimum with the first unsorted element 10 arr[i], arr[min_idx] = arr[min_idx], arr[i] 11 return arr 12 13 # Test 14 numbers = [64, 25, 12, 22, 11] 15 sorted_numbers = selection_sort(numbers) 16 print("Sorted array:", sorted_numbers) 17 18 # Output: 19 # Sorted array: [11, 12, 22, 25, 64] 20</pre>	

	<p>Write a program to execute the insertion sort algorithm.</p> <p>Insertion Sort is a simple sorting algorithm where elements are picked one by one and inserted into their correct position in the sorted portion of the array.</p> <p>"Pick each element and insert it into the sorted part of the array by shifting larger elements to the right."</p>	<pre>main.py > ... 1 def insertion_sort(arr): 2 for i in range(1, len(arr)): 3 key = arr[i] 4 j = i - 1 5 # Move elements greater than key to one position ahead 6 while j >= 0 and arr[j] > key: 7 arr[j + 1] = arr[j] 8 j -= 1 9 arr[j + 1] = key 10 return arr 11 12 # Test 13 numbers = [12, 11, 13, 5, 6] 14 sorted_numbers = insertion_sort(numbers) 15 print("Sorted array:", sorted_numbers) 16 17 # Output: 18 # Sorted array: [5, 6, 11, 12, 13]</pre>	
	<p>Write a program to execute the bubble sort algorithm.</p> <p>Bubble Sort is a simple sorting algorithm where adjacent elements are repeatedly swapped if they are in the wrong order.</p> <p>"Loop through the list multiple times, swapping adjacent elements if they're in the wrong order until sorted."</p>	<pre>main.py > ... 1 def bubble_sort(arr): 2 n = len(arr) 3 for i in range(n): 4 # Last i elements are already sorted 5 for j in range(0, n-i-1): 6 if arr[j] > arr[j+1]: 7 # Swap 8 arr[j], arr[j+1] = arr[j+1], arr[j] 9 return arr 10 11 # Test 12 numbers = [64, 34, 25, 12, 22, 11, 90] 13 sorted_numbers = bubble_sort(numbers) 14 print("Sorted array:", sorted_numbers) 15 16 # Output: 17 # Sorted array: [11, 12, 22, 25, 34, 64, 90] 18</pre>	






	<p>1. Flask là gì? Tại sao nó được gọi là micro-framework?</p> <p>Flask là một micro-framework của Python dùng để xây dựng ứng dụng web một cách nhanh gọn. Nó được gọi là “micro” vì bản thân framework rất tối giản: không tích hợp sẵn ORM, hệ thống xác thực, form handling hay các thành phần nặng khác. Flask chỉ cung cấp những phần cốt lõi như routing, request/response handling và template engine Jinja2, còn mọi tính năng nâng cao đều do lập trình viên tự lựa chọn thông qua các extension như Flask-SQLAlchemy, Flask-Login, Flask-Migrate..., giúp ứng dụng vừa nhẹ, vừa linh hoạt và dễ mở rộng tùy theo nhu cầu.</p>	
	<p>2. Flask định nghĩa route như thế nào, và cơ chế routing hoạt động ra sao?</p> <p>Trong Flask, route được định nghĩa bằng decorator <code>@app.route()</code> đặt lên trên một hàm. Khi Flask khởi động, decorator này đăng ký URL tương ứng vào hệ thống routing của ứng dụng và ánh xạ (map) URL đó đến hàm được gắn decorator. Khi người dùng gửi request đến URL, Flask sẽ duyệt bảng routing, tìm route phù hợp, sau đó gọi hàm xử lý tương ứng (view function), nhận dữ liệu trả về và biến nó thành một HTTP response gửi lại cho client.</p>	
	<p>3. Sự khác nhau giữa request, session, và g trong Flask là gì?</p> <p>request, session và g đều nằm trong hệ thống context của Flask nhưng phục vụ các mục đích khác nhau:</p> <ul style="list-style-type: none"> - request chứa toàn bộ dữ liệu của HTTP request hiện tại, như URL, method, headers, body... Nó chỉ tồn tại trong suốt một request và được dùng để đọc thông tin mà client gửi lên. - session dùng để lưu trữ dữ liệu tạm cho từng người dùng thông qua cookie đã ký (signed cookie). Nó tồn tại qua nhiều request của cùng một user và dùng để lưu trạng thái ngắn hạn như user_id, theme, giỏ hàng... - g (global for one request) là đối tượng tạm để lưu các dữ liệu dùng chung trong một request, ví dụ như kết nối database đã mở sẵn. Nó tồn tại trong vòng đời của request, nhưng không lưu giữ giữa các request. <p>Tóm lại: request để đọc thông tin của HTTP request, session để lưu trạng thái cho user qua nhiều request, còn g để chia sẻ dữ liệu tạm thời trong một request duy nhất.</p>	

<p>Flask</p> <p>4. Giải thích cơ chế "Application Context" và "Request Context" trong Flask.</p> <p>Trong Flask, "context" là cơ chế giúp quản lý các biến và dữ liệu gắn liền với ứng dụng hoặc request mà không cần phải truyền chúng qua từng hàm. Có hai loại context chính:</p> <ol style="list-style-type: none"> 1. Application Context (app context): <ul style="list-style-type: none"> - Chứa thông tin liên quan đến ứng dụng, như current_app và g. - current_app cho phép truy cập app hiện tại mà không cần truyền biến app thủ công. - g là nơi lưu trữ dữ liệu tạm thời dùng chung trong vòng đời của ứng dụng. - Application context tồn tại độc lập với request và có thể được push thủ công khi chạy code ngoài request. 2. Request Context (request context): <ul style="list-style-type: none"> - Chứa thông tin của một HTTP request cụ thể, như request, session. - Flask tự động push request context khi một request đến và pop khi request kết thúc. - Cho phép truy cập các biến như request hay session mà không cần truyền chúng vào từng hàm. <p>Tóm lại: Application Context liên quan đến ứng dụng nói chung, Request Context liên quan đến từng request cụ thể. Cơ chế này giúp Flask duy trì tính thread-safe, cho phép dùng các biến "global" nhưng chỉ gắn với request hoặc app hiện tại, tránh xung đột giữa các request song song.</p>	
<p>5. Làm sao để cấu hình Flask thành production? Dùng WSGI nào?</p> <p>Để chạy Flask trong môi trường production, không nên dùng server phát triển tích hợp (app.run()), vì nó không tối ưu về hiệu năng và bảo mật. Thay vào đó, Flask nên được triển khai thông qua một server WSGI như:</p> <ul style="list-style-type: none"> - Gunicorn: phổ biến, dễ cấu hình, hỗ trợ worker đa luồng và đa tiến trình. - uWSGI: mạnh mẽ, tùy biến cao, hỗ trợ nhiều worker, caching, logging. - Waitress: nhẹ, dễ dùng cho Windows hoặc môi trường đơn giản. <p>Thường, Flask được đặt sau một reverse proxy như Nginx hoặc Apache để xử lý SSL, load balancing, static files, caching và bảo vệ server. Quy trình triển khai production chuẩn thường là: Nginx nhận request → chuyển đến Gunicorn/uWSGI chạy Flask → trả response về client. Ngoài ra, cần cấu hình logging, error handling, limit worker và timeout để đảm bảo ứng dụng ổn định.</p>	

<p>1. FastAPI dựa trên chuẩn ASGI là gì? Khác gì so với WSGI của Flask?</p> <p>ASGI (Asynchronous Server Gateway Interface) là chuẩn giao tiếp giữa server và ứng dụng Python hỗ trợ async I/O, giúp xử lý nhiều kết nối đồng thời một cách hiệu quả. FastAPI dựa trên ASGI, cho phép viết các endpoint bất đồng bộ (async def) để xử lý I/O-bound tasks như gọi API, đọc/ghi database, hay thao tác file mà không block các request khác, từ đó đạt concurrency cao. Ngược lại, Flask dùng WSGI (Web Server Gateway Interface) – synchronous, chỉ xử lý một request tại một thời điểm trong mỗi worker, nên cần nhiều worker để tăng throughput. Khi triển khai FastAPI, thường dùng ASGI server như Uvicorn hoặc Hypercorn để tận dụng async, trong khi Flask thường chạy qua WSGI server như Gunicorn hoặc uWSGI.</p>	
<p>2. FastAPI hỗ trợ khai báo request parameters như thế nào?</p> <p>Trong FastAPI, request parameters có thể được khai báo trực tiếp trong định nghĩa hàm endpoint với kiểu dữ liệu rõ ràng (type hints), và framework sẽ tự động parse và validate:</p> <ul style="list-style-type: none"> - Path parameters: lấy giá trị từ URL. - Query parameters: lấy từ query string. - Request body: dùng Pydantic model để parse JSON body và validate dữ liệu. <p>FastAPI tự động kiểm tra type, bắt lỗi nếu không đúng, và sinh thông báo lỗi chuẩn JSON, giúp giảm code boilerplate và tăng tính an toàn.</p>	
<p>3. Pydantic dùng để làm gì trong FastAPI?</p> <p>Pydantic trong FastAPI được dùng để định nghĩa các model dữ liệu với type hints, giúp:</p> <ol style="list-style-type: none"> 1. Validate dữ liệu đầu vào (data validation): Kiểm tra kiểu dữ liệu, giá trị bắt buộc, định dạng (ví dụ email, URL, số dương...). Nếu dữ liệu không hợp lệ, FastAPI tự động trả lỗi chuẩn JSON. 2. Serialize/deserialize: Chuyển đổi dữ liệu Python → JSON khi trả response, và JSON → Python object khi nhận request body. 3. Tạo schema OpenAPI tự động: FastAPI dùng Pydantic models để sinh tài liệu Swagger / ReDoc, giúp client biết chính xác cấu trúc request/response mà không cần viết thêm code. <p>Tóm lại: Pydantic giúp type-safe, giảm lỗi runtime, và tự động tạo tài liệu API, làm FastAPI rất mạnh trong việc xây dựng API hiện đại và đáng tin cậy.</p>	<pre>◆ main.py > ... 1 @app.get("/items/{item_id}") 2 async def read_item(item_id: int): 3 return {"item_id": item_id} 4 @app.get("/items/") 5 async def read_items(skip: int = 0, limit: int = 10): 6 return {"skip": skip, "limit": limit} 7 8 from pydantic import BaseModel 9 class Item(BaseModel): 10 name: str 11 price: float 12 13 @app.post("/items/") 14 async def create_item(item: Item): 15 return item</pre>

FastAPI	<p>4. FastAPI tự sinh tài liệu API (Swagger / Redoc) như thế nào?</p> <p>FastAPI tự động sinh tài liệu API (Swagger UI và ReDoc) dựa trên thông tin từ type hints và Pydantic models mà bạn định nghĩa trong các endpoint. Khi bạn khai báo các path parameters, query parameters, request body hoặc response model với kiểu dữ liệu cụ thể, FastAPI sẽ:</p> <ol style="list-style-type: none"> 1. Thu thập thông tin type và validation từ Pydantic models và function signature. 2. Sinh ra schema OpenAPI (trước đây là Swagger specification) mô tả các endpoint, request/response, kiểu dữ liệu, giá trị mặc định, và mô tả. 3. Cung cấp giao diện web: /docs → Swagger UI, cho phép test API trực tiếp từ trình duyệt. /redoc → ReDoc, hiển thị tài liệu đẹp, dễ đọc. <p>Điểm mạnh: lập trình viên không cần viết tài liệu thủ công, tất cả được đồng bộ với code, giảm lỗi và luôn cập nhật theo thay đổi của API.</p>	
	<p>5. Sự khác nhau giữa function sync và async trong FastAPI? Khi nào nên dùng async?</p> <p>Trong FastAPI, một endpoint có thể được định nghĩa là synchronous (def) hoặc asynchronous (async def).</p> <ul style="list-style-type: none"> - Sync function (def): Thực thi tuần tự, blocking. Mỗi request phải chờ hàm hoàn tất trước khi worker có thể xử lý request khác. Phù hợp với CPU-bound tasks hoặc các thao tác nhanh, không cần I/O nhiều. - Async function (async def): Sử dụng await để gọi các I/O-bound operations như gọi API, truy vấn database, đọc/ghi file mà không block event loop. Nhiều request có thể được xử lý đồng thời trong cùng một worker, tăng async productivity và high concurrency. <p>Khi nào dùng async:</p> <ul style="list-style-type: none"> - Khi endpoint cần thực hiện nhiều I/O-bound operations chậm (network call, database query). - Khi muốn tối ưu xử lý nhiều request đồng thời mà không tăng số worker. <p>Lưu ý: Async không giúp cải thiện CPU-bound tasks; những tác vụ nặng về CPU vẫn cần xử lý bằng worker riêng hoặc task queue như Celery.</p> <p>Tóm lại: Sử dụng async để tận dụng event loop cho I/O-bound tasks, còn sync phù hợp cho các tác vụ nhanh hoặc CPU-bound.</p>	

Trick	<p>Giải thích:</p> <ul style="list-style-type: none"> - <code>my_list=[]</code> là mutable, được tạo 1 lần khi function định nghĩa. - Các lần gọi sau cùng dùng list cũ, không tạo mới. 	<pre>main.py > ... def add_item(item, my_list=[]): my_list.append(item) return my_list print(add_item(1)) print(add_item(2)) # output # [1] # [1, 2]</pre>	
	<p>Giải thích:</p> <p>Python làm floor division → làm tròn xuống số nguyên gần nhất.</p>	<pre>main.py > ... print(-7 // 3) # output: # -3</pre>	
	<p>Giải thích:</p> <p><code>start:end:step, step=2</code> → lấy 1 ký tự 1 lần nhảy 2.</p>	<pre>main.py > ... a = "Python" print(a[::2]) # output # "Pto"</pre>	
	<p>Giải thích:</p> <ul style="list-style-type: none"> - Python cache số nguyên nhỏ từ -5 đến 256. - Nếu <code>a = 257; b = 257</code> → False. 	<pre>main.py > ... a = 256 b = 256 print(a is b) # output # True</pre>	
	<p>Giải thích:</p> <p>Python cho phép so sánh “chuỗi”: $1 < x \text{ and } x < 10$.</p>	<pre>main.py > ... x = 5 print(1 < x < 10) # output # True</pre>	

	<p>Giải thích:</p> <p>*b lấy tất cả phần “giữa” giữa a và c.</p>		
	<p>Giải thích:</p> <p>tuple immutable, list copy không ảnh hưởng tuple gốc.</p>		
	<p>Giải thích:</p> <p>else chạy khi loop không bị break.</p>		
	<p>Giải thích:</p> <p>Lambda capture biến theo reference, lúc gọi n=2 (lớp cuối cùng). Nếu muốn fix: lambda x, n=n: x+n.</p>		
	<p>Đáp án:</p> <p>Không lỗi, nhưng ghi đè tên đã import, dễ gây nhầm lẫn.</p>		

	<p>Giải thích:</p> <p>Floating point không chính xác tuyệt đối, $0.1+0.2 \approx 0.3000000000000004$.</p>	<pre>main.py > ... 1 print(0.1 + 0.2 == 0.3) 2 3 # output 4 # False 5 </pre>	
	<p>Giải thích:</p> <p>a và b cùng trỏ tới 1 list, thay đổi cùng ảnh hưởng.</p>	<pre>main.py > ... 1 a = b = [] 2 a.append(1) 3 print(b) 4 5 # output 6 # [1] 7 </pre>	
	<p>Giải thích:</p> <ul style="list-style-type: none"> - and trả về giá trị đầu False hoặc giá trị cuối True - or trả về giá trị đầu True hoặc giá trị cuối False 	<pre>main.py > ... 1 x = 0 2 y = 1 3 print(x and y) 4 print(x or y) 5 6 # output 7 # 0 8 # 1 9 </pre>	
	<p>Giải thích:</p> <p>Python intern small strings, tiết kiệm bộ nhớ.</p>	<pre>main.py > ... 1 a = "hello" 2 b = "hello" 3 print(a is b) 4 5 # This will print True because Python caches small strings and reuses them. 6 # output 7 # True 8 </pre>	
	<p>Giải thích:</p> <p>Key trong dict phải immutable, list mutable không dùng được.</p>	<pre>main.py > ... 1 my_dict = {[1,2]: "value"} 2 3 # output 4 # Error: unhashable type: 'list' 5 </pre>	

	<p>What is SQL? Differentiate between DDL, DML, DCL, and TCL.</p> <p>SQL (Structured Query Language) is a standard language used to interact with relational databases. You use SQL to store, retrieve, update, and manage data. Almost all relational database systems—like PostgreSQL, MySQL, SQL Server, Oracle—use SQL as their core language.</p> <p>DDL is used to define or modify the structure of the database, such as creating or altering tables.</p>	<pre>main.sql 1 -- DDL (Data Definition Language) - Ngôn ngữ định nghĩa dữ liệu 2 CREATE TABLE users (3 id SERIAL PRIMARY KEY, 4 name VARCHAR(100) 5); 6 ALTER TABLE users ADD COLUMN email VARCHAR(100); 7 8 -- DML (Data Manipulation Language) - Ngôn ngữ thao tác dữ liệu 9 INSERT INTO users (name, email) VALUES ('Alice', 'alice@example.com'); 10 UPDATE users SET name = 'Alice Nguyen' WHERE id = 1; 11 DELETE FROM users WHERE id = 1; 12 13 -- C. DCL (Data Control Language) - Ngôn ngữ điều khiển dữ liệu 14 GRANT SELECT ON users TO read_only_user; 15 REVOKE SELECT ON users FROM read_only_user; 16 17 -- D. TCL (Transaction Control Language) - Ngôn ngữ điều khiển giao dịch 18 BEGIN; 19 UPDATE accounts SET balance = balance - 100 WHERE id = 1; 20 UPDATE accounts SET balance = balance + 100 WHERE id = 2; 21 COMMIT; -- or ROLLBACK; 22</pre>	
	<p>What is the difference between a Primary Key and a Foreign Key?</p> <p>A Primary Key is a column (or set of columns) that uniquely identifies each row in a table. It guarantees two things:</p> <ol style="list-style-type: none"> 1. Uniqueness – no two rows can have the same primary key value. 2. Non-null – a primary key cannot contain NULL values. <p>A Foreign Key, on the other hand, is a column (or set of columns) in one table that refers to the primary key of another table. Its purpose is to maintain referential integrity by ensuring that the value in the foreign key column exists in the primary key column of the referenced table.</p>	<pre>main.sql 1 -- Parent table 2 CREATE TABLE customers (3 customer_id SERIAL PRIMARY KEY, -- Primary Key 4 name TEXT NOT NULL 5); 6 7 -- Child table 8 CREATE TABLE orders (9 order_id SERIAL PRIMARY KEY, 10 customer_id INT REFERENCES customers(customer_id), -- Foreign Key 11 amount NUMERIC NOT NULL 12); 13</pre>	
	<p>What is an Index? What types of indexes exist?</p> <p>An index in SQL is a database structure that improves the speed of data retrieval operations on a table at the cost of additional storage and slightly slower write operations. Think of it like an index in a book: instead of scanning every page to find a topic, you can go straight to the page number listed in the index.</p> <p>There are several types of indexes in SQL:</p> <ul style="list-style-type: none"> - Primary Key Index: Automatically created when you define a 	<pre>main.sql 1 -- Creating a simple B-Tree index 2 CREATE INDEX idx_employee_name ON employee(name); 3 4 -- Creating a unique index 5 CREATE UNIQUE INDEX idx_employee_email ON employee(email); 6 7 -- Creating a composite index 8 CREATE INDEX idx_employee_name_dept ON employee(name, department_id); 9 10 -- Creating a partial index 11 CREATE INDEX idx_active_employee ON employee(name) WHERE status = 'active'; 12 13 -- Creating an expression index 14 CREATE INDEX idx_lower_email ON employee(LOWER(email));</pre>	
	<p>What is the difference between INNER JOIN, LEFT JOIN, RIGHT JOIN, and FULL OUTER JOIN?</p> <p>In SQL, joins are used to combine rows from two or more tables based on a related column.</p> <ul style="list-style-type: none"> - INNER JOIN returns only the rows where there is a match in both tables. - LEFT JOIN returns all rows from the left table and the matching rows from the right table; if there is no match, the right side will have NULL. 	<pre>main.sql 1 -- INNER JOIN: only matching rows 2 SELECT e.id, e.name, d.name AS department FROM employee e 3 INNER JOIN department d ON e.department_id = d.id; 4 5 -- LEFT JOIN: all employees, even if no department 6 SELECT e.id, e.name, d.name AS department FROM employee e 7 LEFT JOIN department d ON e.department_id = d.id; 8 9 -- RIGHT JOIN: all departments, even if no employee 10 SELECT e.id, e.name, d.name AS department FROM employee e 11 RIGHT JOIN department d ON e.department_id = d.id; 12 13 -- FULL OUTER JOIN: all employees and all departments 14 SELECT e.id, e.name, d.name AS department FROM employee e 15 FULL OUTER JOIN department d ON e.department_id = d.id; 16</pre>	

	<p>What is the difference between WHERE and HAVING</p> <p>In SQL, WHERE is used to filter rows before any groupings are made, while HAVING is used to filter groups after GROUP BY has been applied. In other words, WHERE works on individual rows, and HAVING works on aggregated results like COUNT, SUM, AVG, etc.</p>	<pre>main.sql 1 -- WHERE filters individual rows 2 SELECT * FROM employee 3 WHERE salary > 50000; 4 5 -- HAVING filters groups 6 SELECT department_id, COUNT(*) AS num_employees 7 FROM employee 8 GROUP BY department_id 9 HAVING COUNT(*) > 5; 10 </pre>	
	<p>What is the difference between UNION and UNION ALL?</p> <p>In SQL, UNION combines the result sets of two or more queries and removes duplicate rows, while UNION ALL combines result sets including all duplicates. So, UNION ensures uniqueness but may be slightly slower because it needs to check for duplicates, whereas UNION ALL is faster because it just appends all rows.</p>	<pre>main.sql 1 -- Using UNION (duplicates removed) 2 SELECT name FROM employee_2023 3 UNION 4 SELECT name FROM employee_2024; 5 6 -- Using UNION ALL (duplicates kept) 7 SELECT name FROM employee_2023 8 UNION ALL 9 SELECT name FROM employee_2024; 10 </pre>	
	<p>What is the difference between TRUNCATE, DELETE, and DROP?</p> <p>In SQL, DELETE, TRUNCATE, and DROP are used to remove data, but they work differently.</p> <ul style="list-style-type: none"> - DELETE removes specific rows from a table based on a condition and can be rolled back if inside a transaction. - TRUNCATE removes all rows from a table quickly, resets identity counters, but usually cannot be rolled back in many databases. 	<pre>main.sql 1 -- DELETE: remove specific rows 2 DELETE FROM employee WHERE department_id = 5; 3 4 -- TRUNCATE: remove all rows quickly 5 TRUNCATE TABLE employee; 6 7 -- DROP: remove the entire table 8 DROP TABLE employee; 9 </pre>	
	<p>What is the difference between CHAR and VARCHAR?</p> <p>In SQL, CHAR and VARCHAR are used to store text, but they work differently. CHAR(n) is a fixed-length string, meaning it always uses n characters; if the input is shorter, it pads with spaces. VARCHAR(n) is a variable-length string, storing only the characters you input up to n, so it's more space-efficient. Use CHAR when the length of data is consistent, and VARCHAR when it varies.</p>	<pre>main.sql 1 -- CHAR: fixed length of 10 2 CREATE TABLE employee_char (3 name CHAR(10) 4); 5 6 -- VARCHAR: variable length up to 50 7 CREATE TABLE employee_varchar (8 name VARCHAR(50) 9); 10 11 -- Example inserts 12 INSERT INTO employee_char (name) VALUES ('Tom'); -- stored as 'Tom' 13 INSERT INTO employee_varchar (name) VALUES ('Tom'); -- stored as 'Tom' 14 </pre>	

	<p>What is SQL Injection? How can it be prevented?</p> <p>SQL Injection is a security vulnerability that occurs when an attacker can manipulate an SQL query by inserting malicious input, potentially allowing them to read, modify, or delete data in the database. It happens when user input is directly concatenated into SQL statements without proper validation.</p> <p>Prevention methods:</p> <ul style="list-style-type: none"> - Use parameterized queries / prepared statements: never 	<pre>main.sql 1 -- Unsafe: vulnerable to SQL Injection 2 -- user_input = '1 OR 1=1' 3 SELECT * FROM employee WHERE id = 'user_input'; 4 5 -- Safe: using parameterized query 6 PREPARE get_employee(int) AS 7 SELECT * FROM employee WHERE id = \$1; 8 9 EXECUTE get_employee(1);</pre>	
	<p>Write a query to find the maximum/minimum value in a column.</p> <p>In SQL, you can use the MAX() function to find the largest value in a column, and MIN() to find the smallest value. These are aggregate functions that scan the column and return a single value.</p>	<pre>main.sql 1 -- Find the maximum salary 2 SELECT MAX(salary) AS max_salary 3 FROM employee; 4 5 -- Find the minimum salary 6 SELECT MIN(salary) AS min_salary 7 FROM employee;</pre>	
	<p>Write a query to count records by group.</p> <p>In SQL, to count records by group, you use the COUNT() aggregate function together with GROUP BY. This allows you to count how many rows belong to each category or group in a column.</p>	<pre>main.sql 1 -- Count departments with more than 5 employees 2 SELECT department_id, COUNT(*) AS num_employees 3 FROM employee 4 GROUP BY department_id 5 HAVING COUNT(*) > 5;</pre>	
	<p>How do you retrieve the 5 most recent records from a table?</p> <p>In SQL, to retrieve the most recent records, you typically need a column that indicates order, like a date, timestamp, or an auto-incrementing ID. You use ORDER BY to sort the records in descending order and LIMIT to restrict the number of rows returned.</p>	<pre>main.sql 1 -- Get 5 most recent employees by created_date 2 SELECT * FROM employee 3 ORDER BY created_date DESC 4 LIMIT 5; 5 6 -- Get 5 most recent employees by auto-increment ID 7 SELECT * FROM employee 8 ORDER BY id DESC 9 LIMIT 5;</pre>	
	<p>Write a query to filter data based on multiple conditions.</p> <p>In SQL, to filter data based on multiple conditions, you use the WHERE clause with AND, OR, and parentheses () to combine conditions logically. This allows you to precisely select rows that meet complex criteria.</p>	<pre>main.sql 1 -- Filter employees in department 1 with salary > 50000 2 SELECT * FROM employee 3 WHERE department_id = 1 4 AND salary > 50000; 5 6 -- Filter employees in department 1 or department 2 with salary > 50000 7 SELECT * FROM employee 8 WHERE (department_id = 1 OR department_id = 2) 9 AND salary > 50000;</pre>	

	<p>How do you combine data from multiple tables?</p> <p>In SQL, to combine data from multiple tables, you use JOIN clauses. The most common types are INNER JOIN, LEFT JOIN, RIGHT JOIN, and FULL OUTER JOIN. JOINs allow you to match rows from one table with rows from another table based on a related column, usually a primary key and foreign key relationship.</p>	<pre>main.sql 1 -- Combine employee data with department data using INNER JOIN 2 SELECT e.id, e.name, d.name AS department 3 FROM employee e 4 INNER JOIN department d ON e.department_id = d.id; 5 6 -- LEFT JOIN example: get all employees, including those without a department 7 SELECT e.id, e.name, d.name AS department 8 FROM employee e 9 LEFT JOIN department d ON e.department_id = d.id; 10</pre>	
	<p>Write a query to find records that appear more than once.</p> <p>In SQL, to find records that appear more than once, you use GROUP BY on the column(s) you want to check and HAVING COUNT(*) > 1 to filter groups with duplicates.</p>	<pre>main.sql 1 -- Find duplicate employee names 2 SELECT name, COUNT(*) AS count_name 3 FROM employee 4 GROUP BY name 5 HAVING COUNT(*) > 1; 6 7 -- Find duplicate emails 8 SELECT email, COUNT(*) AS count_email 9 FROM employee 10 GROUP BY email 11 HAVING COUNT(*) > 1; 12</pre>	
	<p>Write a query to update data in one table based on data from another table.</p> <p>In SQL, to update data in one table based on data from another table, you use the UPDATE statement combined with a JOIN (or a subquery) to match rows between the tables. This lets you change values in one table according to corresponding values in another table.</p>	<pre>main.sql 1 -- Update employee department names based on department table 2 UPDATE employee e 3 SET department_name = d.name 4 FROM department d 5 WHERE e.department_id = d.id; 6 7 -- Another way using subquery 8 UPDATE employee 9 SET department_name = (10 SELECT name 11 FROM department 12 WHERE department.id = employee.department_id 13); 14</pre>	
	<p>What is the difference between NULL and an empty string?</p> <p>In SQL, NULL represents the absence of a value — it means the value is unknown or missing. An empty string (""), on the other hand, is a known value that just happens to have no characters. NULL and empty string are not the same, and comparisons with NULL require special handling using IS NULL or IS NOT NULL.</p>	<pre>main.sql 1 -- NULL example 2 SELECT * FROM employee 3 WHERE middle_name IS NULL; -- selects employees with unknown middle name 4 5 -- Empty string example 6 SELECT * FROM employee 7 WHERE middle_name = ''; -- selects employees whose middle name is explicitly empty 8 9 -- Optional: checking differences / Kiểm tra sự khác biệt 10 SELECT 11 COUNT(*) AS null_count, 12 COUNT(CASE WHEN middle_name = '' THEN 1 END) AS empty_string_count 13 FROM employee; 14</pre>	
	<p>How do you sort query results in ascending and descending order?</p> <p>In SQL, to sort query results, you use the ORDER BY clause. By default, ORDER BY sorts in ascending order (ASC), and you can explicitly use DESC for descending order. You can also sort by multiple columns, each with its own order.</p>	<pre>main.sql 1 -- Sort employees by salary ascending 2 SELECT * FROM employee 3 ORDER BY salary ASC; 4 5 -- Sort employees by salary descending 6 SELECT * FROM employee 7 ORDER BY salary DESC; 8 9 -- Sort by department ascending and then salary descending 10 SELECT * FROM employee 11 ORDER BY department_id ASC, salary DESC; 12</pre>	
	<p>What are aggregate functions in SQL? Give examples.</p> <p>In SQL, aggregate functions are functions that perform calculations on a set of values and return a single summary value. They are often used with GROUP BY to summarize data by groups. Common aggregate functions include COUNT(), SUM(), AVG(), MAX(), MIN().</p>	<pre>main.sql 1 -- Count the total number of employees 2 SELECT COUNT(*) AS total_employees FROM employee; 3 4 -- Sum of all salaries 5 SELECT SUM(salary) AS total_salary FROM employee; 6 7 -- Average salary 8 SELECT AVG(salary) AS average_salary FROM employee; 9 10 -- Maximum and minimum salary 11 SELECT MAX(salary) AS max_salary, MIN(salary) AS min_salary FROM employee; 12 13 -- Count employees in each department 14 SELECT department_id, COUNT(*) AS num_employees FROM employee 15 GROUP BY department_id; 16</pre>	

	<p>What is the difference between BETWEEN and IN operators?</p> <p>In SQL, BETWEEN is used to filter values within a range, including the boundary values. IN is used to filter values that match any value in a list. BETWEEN is best for continuous ranges, while IN is best for discrete sets of values.</p>	<pre>main.sql 1 -- Using BETWEEN: select employees with salary between 40000 and 60000 2 SELECT * FROM employee WHERE salary BETWEEN 40000 AND 60000; 3 4 -- Using IN: select employees in departments 1, 2, or 3 5 SELECT * FROM employee WHERE department_id IN (1, 2, 3); 6 7 -- Optional: combine with NOT / Kết hợp với NOT 8 -- NOT BETWEEN 9 SELECT * FROM employee WHERE salary NOT BETWEEN 40000 AND 60000; 10 11 -- NOT IN 12 SELECT * FROM employee WHERE department_id NOT IN (1, 2, 3); 13 </pre>	
	<p>How do you find the length of a string or number of characters in a column?</p> <p>In SQL, to find the length of a string or the number of characters in a column, you use the LENGTH() function in PostgreSQL. This function returns the total number of characters, including spaces.</p>	<pre>main.sql 1 -- Find the length of employee names 2 SELECT name, LENGTH(name) AS name_length FROM employee; 3 4 -- Example: find employees with name longer than 10 characters 5 SELECT * FROM employee WHERE LENGTH(name) > 10; 6 7 -- Optional: using for numeric values / Dùng cho số: 8 -- Convert number to text to find length 9 SELECT id, LENGTH(CAST(id AS TEXT)) AS id_length FROM employee; 10 </pre>	
	<p>What is the difference between DISTINCT and no DISTINCT in a SELECT statement?</p> <p>In SQL, DISTINCT is used in a SELECT statement to return only unique values, removing duplicates from the result set. Without DISTINCT, the query returns all matching rows, including duplicates. Use DISTINCT when you want to avoid repeated values in your results.</p>	<pre>main.sql 1 -- Without DISTINCT: may return duplicate department IDs 2 SELECT department_id FROM employee; 3 4 -- With DISTINCT: only unique department IDs 5 SELECT DISTINCT department_id FROM employee; 6 7 -- Example with multiple columns 8 SELECT DISTINCT department_id, job_title FROM employee; 9 </pre>	
	<p>How do you rename a column or table in a query result?</p> <p>In SQL, to rename a column or table in the query result, you use aliases with the AS keyword. Aliases do not change the actual table or column names in the database; they only rename them in the result set for readability.</p>	<pre>main.sql 1 -- Rename a column in the result 2 SELECT name AS employee_name, salary AS employee_salary FROM employee; 3 4 -- Rename a table in the query 5 SELECT e.name, e.salary FROM employee AS e; 6 7 -- Using both column and table aliases 8 SELECT e.name AS employee_name, e.salary AS employee_salary FROM employee AS e; 9 </pre>	
	<p>How do you limit the number of rows returned in a query?</p> <p>In SQL, to limit the number of rows returned by a query, you use the LIMIT clause. This allows you to retrieve only a specific number of rows instead of the entire result set. In PostgreSQL, you can also combine it with OFFSET to skip a certain number of rows.</p>	<pre>main.sql 1 -- Get only 5 rows 2 SELECT * FROM employee 3 LIMIT 5; 4 5 -- Get 5 rows starting from the 6th row 6 SELECT * FROM employee 7 LIMIT 5 OFFSET 5; 8 </pre>	
	<p>What is the difference between a PRIMARY KEY and a UNIQUE constraint?</p> <p>In SQL, both PRIMARY KEY and UNIQUE constraints enforce uniqueness for a column or a set of columns. The main differences are: a PRIMARY KEY uniquely identifies each row in a table and cannot accept NULL values, and a table can have only one primary key. A UNIQUE constraint also ensures uniqueness but can accept one NULL value per column in most databases, and a table can have multiple UNIQUE constraints.</p>	<pre>main.sql 1 -- PRIMARY KEY 2 CREATE TABLE employee (3 id SERIAL PRIMARY KEY, 4 name VARCHAR(50) 5); 6 7 -- UNIQUE constraint 8 CREATE TABLE employee_email (9 id SERIAL PRIMARY KEY, 10 email VARCHAR(100) UNIQUE 11); 12 13 -- Optional: multiple UNIQUE constraints 14 CREATE TABLE example (15 col1 INT UNIQUE, 16 col2 INT UNIQUE 17); 18 </pre>	

	<p>What is the difference between CHECK and DEFAULT constraints?</p> <p>In SQL, CHECK and DEFAULT constraints serve different purposes. A CHECK constraint enforces a rule on the values that can be inserted into a column, ensuring data validity. A DEFAULT constraint automatically assigns a default value to a column when no value is provided during insertion.</p>	<pre>main.sql 1 -- CHECK constraint: salary must be positive 2 CREATE TABLE employee (3 id SERIAL PRIMARY KEY, 4 name VARCHAR(50), 5 salary NUMERIC CHECK (salary > 0) 6); 7 8 -- DEFAULT constraint: department defaults to 'General' 9 CREATE TABLE employee_default (10 id SERIAL PRIMARY KEY, 11 name VARCHAR(50), 12 department VARCHAR(50) DEFAULT 'General' 13); 14 15 -- Example insert 16 INSERT INTO employee_default (name) VALUES ('Tom'); 17 -- department will automatically be 'General'</pre>	
	<p>What is a cascade delete? When should it be used?</p> <p>In SQL, a cascade delete is a foreign key option that automatically deletes child rows when the corresponding parent row is deleted. This ensures referential integrity without leaving orphaned records in related tables.</p> <p>You should use cascade delete when you want related records to be automatically removed together with the parent, such as deleting an order and all its order items. Be careful using it, because it can delete large amounts of data unintentionally if not planned properly.</p>	<pre>main.sql 1 -- Parent table: department 2 CREATE TABLE department (3 id SERIAL PRIMARY KEY, 4 name VARCHAR(50) 5); 6 7 -- Child table: employee with foreign key and cascade delete 8 CREATE TABLE employee (9 id SERIAL PRIMARY KEY, 10 name VARCHAR(50), 11 department_id INT REFERENCES department(id) ON DELETE CASCADE 12); 13 14 -- Deleting a department automatically deletes all employees in that department 15 DELETE FROM department WHERE id = 1; 16 </pre>	
	<p>What is the difference between ON DELETE CASCADE and ON DELETE SET NULL?</p> <p>In SQL, both ON DELETE CASCADE and ON DELETE SET NULL are foreign key actions triggered when a parent row is deleted. ON DELETE CASCADE automatically deletes all child rows that reference the parent, ensuring no orphaned records remain. ON DELETE SET NULL does not delete child rows; instead, it sets the foreign key column in the child table to NULL, keeping the child rows but removing the reference to the deleted parent.</p>	<pre>main.sql 1 -- Parent table: department 2 CREATE TABLE department (3 id SERIAL PRIMARY KEY, 4 name VARCHAR(50) 5); 6 7 -- Child table with ON DELETE CASCADE 8 CREATE TABLE employee_cascade (9 id SERIAL PRIMARY KEY, 10 name VARCHAR(50), 11 department_id INT REFERENCES department(id) ON DELETE CASCADE 12); 13 14 -- Child table with ON DELETE SET NULL 15 CREATE TABLE employee_setnull (16 id SERIAL PRIMARY KEY, 17 name VARCHAR(50), 18 department_id INT REFERENCES department(id) ON DELETE SET NULL 19); 20 21 -- Delete department with id = 1 22 DELETE FROM department WHERE id = 1; 23 24 -- In employee_cascade: all employees with department_id = 1 are deleted 25 -- In employee_setnull: all employees with department_id = 1 have department_id set to NULL 26 </pre>	

	<p>How can you optimize a slow query?</p> <p>In SQL, to optimize a slow query, you can use several strategies. First, analyze and add indexes on columns used in WHERE, JOIN, ORDER BY, or GROUP BY clauses to speed up lookups. Second, **avoid SELECT *** and only retrieve the columns you need. Third, rewrite complex queries using joins, subqueries, or CTEs more efficiently. Fourth, check execution plans with EXPLAIN to find bottlenecks. Finally, consider caching, partitioning, or denormalization if appropriate.</p> <p>Optional tips:</p> <ul style="list-style-type: none"> - Ensure indexes are used on filtering and joining columns. - Avoid functions on indexed columns in WHERE clauses (e.g., LENGTH(name) > 5 may bypass index). - Limit data scanned with proper WHERE conditions. 	<pre>main.sql 1 -- Check how the query runs 2 EXPLAIN ANALYZE 3 SELECT e.id, e.name, d.name AS department 4 FROM employee e 5 JOIN department d ON e.department_id = d.id 6 WHERE e.salary > 50000; 7</pre>
	<p>When should you use an Index? When should you avoid using one?</p> <p>In SQL, you should use an index when you have frequent queries that filter, sort, or join on a column, especially for large tables. Indexes can drastically speed up SELECT queries by allowing the database to find rows without scanning the entire table.</p> <p>You should avoid using an index when: the table is small (full scan is fast enough), the column is updated very frequently (indexes slow down INSERT/UPDATE/DELETE), or there are too many indexes that increase storage and maintenance overhead.</p>	<pre>main.sql 1 -- Create an index on the salary column 2 CREATE INDEX idx_employee_salary ON employee(salary); 3 4 -- Use index when filtering by salary 5 SELECT * 6 FROM employee 7 WHERE salary > 50000; 8 9 10 -- Optional: Avoid indexing small or frequently updated columns 11 -- Not recommended for a column updated every row frequently 12 CREATE INDEX idx_temp_data ON employee(temp_value); 13</pre>
	<p>What is the difference between a clustered and a non-clustered index?</p> <p>In SQL, a clustered index determines the physical order of data in a table. There can be only one clustered index per table, and the table rows are stored in the order of the clustered index. A non-clustered index does not change the physical order of the table; it creates a separate structure pointing to the actual data, and a table can have multiple non-clustered indexes. Clustered indexes are good for range queries, while non-clustered indexes are good for lookups on specific columns.</p> <p>Key points:</p> <ul style="list-style-type: none"> - Clustered index = physical order of rows. - Non-clustered index = separate structure pointing to rows. - One table → 1 clustered index, many non-clustered indexes possible. 	<pre>main.sql 1 -- PostgreSQL automatically creates a clustered index for PRIMARY KEY 2 CREATE TABLE employee (3 id SERIAL PRIMARY KEY, 4 name VARCHAR(50), 5 salary NUMERIC 6); 7 8 -- Create a non-clustered index (PostgreSQL uses B-tree by default) 9 CREATE INDEX idx_employee_salary ON employee(salary); 10 11 -- Optional: cluster table manually on a column (PostgreSQL) 12 CLUSTER employee USING idx_employee_salary; 13</pre>

	<p>What is an Explain Plan? What is it used for?</p> <p>In SQL, an Explain Plan is a tool that shows how the database executes a query, including the steps, order of operations, and indexes used. It is mainly used to analyze and optimize queries by identifying bottlenecks or inefficient operations.</p> <p>Optional:</p> <ul style="list-style-type: none"> - Helps see which indexes are used - Helps estimate query cost - Helps identify full table scans 	<pre>main.sql 1 -- Check the execution plan of a query 2 EXPLAIN 3 SELECT e.id, e.name, d.name AS department 4 FROM employee e 5 JOIN department d ON e.department_id = d.id 6 WHERE e.salary > 50000; 7 8 -- Check the actual execution with timing 9 EXPLAIN ANALYZE 10 SELECT e.id, e.name, d.name AS department 11 FROM employee e 12 JOIN department d ON e.department_id = d.id 13 WHERE e.salary > 50000; 14 </pre>	
	<p>What is the difference between a transaction and a lock?</p> <p>In SQL, a transaction is a logical unit of work that consists of one or more SQL statements executed as a single unit. Transactions follow the ACID properties (Atomicity, Consistency, Isolation, Durability) to ensure data integrity. A lock is a mechanism used by the database to control concurrent access to data, preventing conflicts like dirty reads or lost updates. While transactions define what work should be done atomically, locks are the tools the database uses internally to manage concurrency during transactions.</p> <p>Quick notes:</p> <ul style="list-style-type: none"> - Transaction = logical unit of work - Lock = mechanism to control concurrent access - Transactions use locks internally to protect data 	<pre>main.sql 1 -- Start a transaction 2 BEGIN; 3 4 -- Update a row 5 UPDATE employee 6 SET salary = salary + 5000 7 WHERE id = 1; 8 9 -- Commit or rollback 10 COMMIT; -- apply changes 11 -- ROLLBACK; -- undo changes 12 13 -- Example of a lock on a table 14 LOCK TABLE employee IN ACCESS EXCLUSIVE MODE; 15 </pre>	
	<p>Explain the difference between INNER JOIN and CROSS JOIN.</p> <p>An INNER JOIN returns rows from two tables where there is a match based on a specified condition, usually involving primary and foreign keys. Only the rows satisfying the condition appear in the result.</p> <p>A CROSS JOIN, on the other hand, returns the Cartesian product of the two tables, meaning every row from the first table is combined with every row from the second table, regardless of any condition. CROSS JOIN can produce a very large number of rows if both tables are big.</p> <p>Quick notes:</p> <ul style="list-style-type: none"> - INNER JOIN = match rows based on condition - CROSS JOIN = all combinations of rows - Use CROSS JOIN carefully, may create large result sets 	<pre>main.sql 1 -- INNER JOIN example 2 SELECT e.name, d.name AS department 3 FROM employee e 4 INNER JOIN department d ON e.department_id = d.id; 5 6 -- CROSS JOIN example 7 SELECT e.name, d.name AS department 8 FROM employee e 9 CROSS JOIN department d; 10 </pre>	

	<p>What is a subquery and when would you use one?</p> <p>A subquery is a query nested inside another query, typically inside a SELECT, INSERT, UPDATE, or DELETE statement. Subqueries are used when you need to retrieve a value or set of values from one table to use in another query. They help break complex queries into manageable parts and can be used for filtering, calculating aggregates, or checking existence.</p> <p>Quick notes: Subquery = a query inside another query Can return single value, list, or table Useful for filtering, aggregating, or conditional checks</p>	<pre>main.sql 1 -- Subquery in WHERE clause 2 SELECT name, salary 3 FROM employee 4 WHERE department_id = (5 SELECT id 6 FROM department 7 WHERE name = 'Sales' 8); 9 10 -- Subquery in SELECT clause 11 SELECT name, (SELECT name FROM department d WHERE d.id = e.department_id) AS department_name 12 FROM employee e; 13</pre>
	<p>What is a correlated subquery? Give an example.</p> <p>A correlated subquery is a subquery that depends on a column from the outer query. It is executed once for each row of the outer query, unlike a regular subquery which runs independently. Correlated subqueries are useful for row-by-row comparisons or calculations that reference the outer query's data.</p> <p>Quick notes: Correlated subquery references outer query columns Executed once per row of the outer query Useful for comparisons and conditional calculations</p>	<pre>main.sql 1 -- Find employees whose salary is greater than the average salary of their department 2 SELECT e1.name, e1.salary, e1.department_id 3 FROM employee e1 4 WHERE e1.salary > (5 SELECT AVG(e2.salary) 6 FROM employee e2 7 WHERE e2.department_id = e1.department_id 8); 9</pre>
	<p>Explain the difference between EXISTS and IN.</p> <p>In SQL, both EXISTS and IN are used to filter rows based on another query, but they work differently. EXISTS checks if a subquery returns any rows and returns TRUE or FALSE for each row of the outer query; it is often more efficient with large datasets. IN checks if a column's value matches any value in a list or subquery. EXISTS is typically faster when the subquery returns many rows, while IN is more readable for small lists.</p> <p>Quick notes: EXISTS = TRUE/FALSE based on subquery result IN = matches values in list or subquery EXISTS often faster with large subqueries</p>	<pre>main.sql 1 -- Using IN 2 SELECT name FROM employee 3 WHERE department_id IN (4 SELECT id 5 FROM department 6 WHERE name LIKE 'S%' 7); 8 9 -- Using EXISTS 10 SELECT name FROM employee e 11 WHERE EXISTS (12 SELECT 1 13 FROM department d 14 WHERE d.id = e.department_id 15 AND d.name LIKE 'S%' 16); 17</pre>

	<p>How can you enforce data integrity using constraints other than PRIMARY KEY and FOREIGN KEY?</p> <p>In SQL, besides PRIMARY KEY and FOREIGN KEY, you can enforce data integrity using several other constraints: UNIQUE, CHECK, and DEFAULT. UNIQUE ensures that a column or combination of columns has distinct values. CHECK enforces a rule on column values, such as limiting salaries to be positive. DEFAULT automatically assigns a default value to a column if none is provided. These constraints help maintain valid, consistent, and predictable data in the database.</p>	<pre>main.sql 1 CREATE TABLE employee (2 id SERIAL PRIMARY KEY, 3 name VARCHAR(50) UNIQUE, -- Ràng buộc UNIQUE 4 salary NUMERIC CHECK (salary > 0), -- Ràng buộc CHECK 5 department VARCHAR(50) DEFAULT 'General' -- Ràng buộc DEFAULT 6); 7 </pre>
	<p>What is the difference between CASE and IF in SQL?</p> <p>In SQL, CASE and IF are used for conditional logic, but they are different. CASE is part of the SQL standard and works in SELECT, UPDATE, ORDER BY, and other statements; it can evaluate multiple conditions and return different results based on them. IF is usually a procedural statement used in database programming languages like PL/pgSQL, T-SQL, or inside stored procedures, not in standard SELECT queries.</p>	<pre>main.sql 1 -- Using CASE in a SELECT query 2 SELECT name, salary, 3 CASE 4 WHEN salary > 50000 THEN 'High' 5 WHEN salary BETWEEN 30000 AND 50000 THEN 'Medium' 6 ELSE 'Low' 7 END AS salary_level 8 FROM employee; 9 </pre>
	<p>What is a composite key? How is it different from a single-column key?</p> <p>A composite key is a primary key made up of two or more columns in a table, where the combination of values in these columns uniquely identifies each row. A single-column key (or simple primary key) uses only one column to uniquely identify rows. Composite keys are used when no single column can guarantee uniqueness, but a combination of columns can.</p>	<pre>main.sql 1 -- Ví dụ composite key 2 CREATE TABLE order_item (3 order_id INT, 4 product_id INT, 5 quantity INT, 6 PRIMARY KEY (order_id, product_id) -- khóa composite 7); 8 9 -- Ví dụ khóa một cột 10 CREATE TABLE employee (11 id SERIAL PRIMARY KEY, -- khóa một cột 12 name VARCHAR(50) 13); 14 </pre>

	<p>What are transactions in SQL, and what are the ACID properties?</p> <p>In SQL, a transaction is a sequence of operations that are treated as a single logical unit of work, meaning either all operations succeed or none of them take effect. Transactions ensure data consistency, especially in cases of errors or system failures. The reliability of transactions is defined by the ACID properties:</p> <ul style="list-style-type: none"> - Atomicity: all operations in the transaction must succeed or all must be rolled back. - Consistency: the database must move from one valid state to another. - Isolation: transactions should not interfere with each other; intermediate states are not visible. - Durability: once a transaction is committed, the changes are permanently saved even if the system crashes. 	<pre>main.sql 1 BEGIN; 2 3 UPDATE accounts SET balance = balance - 100 WHERE id = 1; 4 UPDATE accounts SET balance = balance + 100 WHERE id = 2; 5 6 COMMIT; -- or ROLLBACK; 7</pre>
	<p>How can you prevent deadlocks in a database?</p> <p>To prevent deadlocks in a database, you should follow good transaction and locking practices. The most important strategy is keeping transactions short and consistent so locks are held for minimal time. Always access tables and rows in a consistent order across all parts of the application to avoid circular waits. Use proper indexing to reduce full table scans that can escalate locks. Avoid unnecessary locking by selecting only the rows you need, and use lower isolation levels when possible to reduce locking intensity. Finally, monitor queries and use tools like EXPLAIN to detect slow operations that may cause lock contention.</p>	<pre>main.sql 1 -- Truy cập bảng theo thứ tự nhất quán 2 BEGIN; 3 UPDATE orders SET status = 'paid' WHERE id = 10; 4 UPDATE order_items SET shipped = true WHERE order_id = 10; 5 6 COMMIT;</pre>
	<p>What is the difference between ROW_NUMBER(), RANK(), and DENSE_RANK()?</p> <p>ROW_NUMBER(), RANK(), and DENSE_RANK() are window functions used to assign ordering numbers to rows, but they behave differently when duplicates appear. ROW_NUMBER() assigns a unique sequential number to each row, even if values are tied. RANK() assigns the same rank to tied rows but skips numbers afterward. DENSE_RANK() also assigns the same rank to tied rows but does not skip numbers, making the sequence continuous.</p> <p>Quick notes:</p> <ul style="list-style-type: none"> - ROW_NUMBER(): no ties, always 1,2,3,... - RANK(): ties share rank, numbers are skipped - DENSE_RANK(): ties share rank, no skipping 	<pre>main.sql 1 SELECT name, score, 2 ROW_NUMBER() OVER (ORDER BY score DESC) AS row_num, 3 RANK() OVER (ORDER BY score DESC) AS rnk, 4 DENSE_RANK() OVER (ORDER BY score DESC) AS dense_rnk 5 6 FROM students;</pre>

	<p>What is the difference between INNER JOIN and EXISTS?</p> <p>The difference between INNER JOIN and EXISTS lies in how they filter data and how they perform. INNER JOIN combines rows from two tables based on matching values; it returns the actual columns from both tables. It is best when you need data from both sides. EXISTS, on the other hand, checks whether a subquery returns at least one row; it returns TRUE or FALSE for each row of the outer query and does not return columns from the subquery. EXISTS is usually more efficient when you only need to check the existence of related rows, especially when the subquery can stop early once a match is found.</p> <p>Quick notes:</p> <ul style="list-style-type: none"> - INNER JOIN returns combined data from two tables. - EXISTS only checks for existence, no need to return columns. - EXISTS can be faster when the subquery has many rows. 	<pre>main.sql 1 -- Ví dụ INNER JOIN 2 SELECT e.name, d.name AS department 3 FROM employee e 4 INNER JOIN department d ON e.department_id = d.id; 5 6 -- Ví dụ EXISTS 7 SELECT e.name 8 FROM employee e 9 WHERE EXISTS (10 SELECT 1 11 FROM department d 12 WHERE d.id = e.department_id 13); 14</pre>
	<p>Write a query to paginate data.</p> <p>A common way to paginate data in SQL is to use the keywords LIMIT and OFFSET. LIMIT controls how many rows you want to return, while OFFSET tells the query how many rows to skip before starting to return results. This is useful when you want to display data page by page, for example in a web application. For page-based pagination, you usually calculate the offset as: (page_number - 1) * page_size.</p>	<pre>main.sql 1 SELECT * 2 FROM employees 3 ORDER BY id 4 LIMIT 10 OFFSET 20; 5 -- Truy vấn này sẽ trả về trang số 3 nếu mỗi trang có 10 dòng, vì nó bỏ qua 20 dòng đầu và lấy 10 dòng tiếp theo.</pre>
	<p>How do you handle NULL values in SQL?</p> <p>In SQL, NULL represents missing or unknown data. Handling NULLs properly is important to avoid errors or incorrect results. You can use:</p> <ul style="list-style-type: none"> - IS NULL / IS NOT NULL to filter rows with or without NULL values. - COALESCE() to replace NULL with a default value. - NULLIF() to return NULL if two expressions are equal. - Aggregate functions like COUNT(column) ignore NULLs, but COUNT(*) counts all rows. <p>Handling NULLs carefully ensures correct calculations and filtering.</p>	<pre>main.sql 1 -- Lọc các dòng có salary là NULL 2 SELECT * FROM employees WHERE salary IS NULL; 3 4 -- Thay thế NULL bằng giá trị mặc định 5 SELECT name, COALESCE(salary, 0) AS salary FROM employees; 6 7 -- Trả về NULL nếu hai giá trị bằng nhau 8 SELECT name, NULLIF(salary, 0) AS salary_or_null FROM employees; 9</pre>

	<p>What is the difference between a STORED PROCEDURE and a FUNCTION?</p> <p>In SQL, both stored procedures and functions are reusable blocks of code stored in the database, but they have key differences. A stored procedure can perform operations like INSERT, UPDATE, DELETE, or complex logic, and it does not have to return a value. It is called with the CALL statement. A function must return a value and can be used in SQL expressions, like in SELECT or WHERE clauses. Functions are generally used for calculations or transformations, while procedures are used for executing business logic or workflows.</p>	<pre>main.sql 1 -- Ví dụ function 2 CREATE FUNCTION get_employee_salary(emp_id INT) 3 RETURNS NUMERIC AS \$\$ 4 BEGIN 5 RETURN (SELECT salary FROM employee WHERE id = emp_id); 6 END; 7 \$\$ LANGUAGE plpgsql; 8 9 -- Ví dụ stored procedure 10 CREATE PROCEDURE raise_salary(emp_id INT, amount NUMERIC) 11 LANGUAGE plpgsql 12 AS \$\$ 13 BEGIN 14 UPDATE employee 15 SET salary = salary + amount 16 WHERE id = emp_id; 17 END; 18 \$\$;</pre>	
	<p>What is the difference between a VIEW and a MATERIALIZED VIEW?</p> <p>A VIEW is a virtual table in SQL that is defined by a query. It does not store data physically, so each time you query it, the database executes the underlying query. A MATERIALIZED VIEW, on the other hand, stores the query results physically. This can improve performance for complex queries because the data is precomputed, but it may become stale if the underlying tables change, and may require manual refresh.</p> <p>Quick notes:</p> <ul style="list-style-type: none"> - VIEW = virtual table, no physical storage - MATERIALIZED VIEW = stores results physically - Use materialized view for performance on heavy queries 	<pre>main.sql 1 -- Tạo view thông thường 2 CREATE VIEW employee_view AS 3 SELECT id, name, salary 4 FROM employee 5 WHERE salary > 50000; 6 7 -- Tạo materialized view 8 CREATE MATERIALIZED VIEW employee_mat_view AS 9 SELECT id, name, salary 10 FROM employee 11 WHERE salary > 50000; 12 13 -- Làm mới materialized view 14 REFRESH MATERIALIZED VIEW employee_mat_view; 15 </pre>	
	<p>How do you find products that have not been included in any orders?</p> <p>To find products that have not been included in any orders, you can use either a LEFT JOIN with a NULL check or a NOT EXISTS subquery. The LEFT JOIN approach joins products with orders and checks for rows where the order does not exist. The NOT EXISTS approach checks for products where no matching order is found. Both methods work, but NOT EXISTS is often more readable for this case.</p> <p>Quick notes:</p> <ul style="list-style-type: none"> - LEFT JOIN + NULL = find unmatched rows - NOT EXISTS = check absence of related rows - Both achieve the same result; choice depends on readability and query planner 	<pre>main.sql 1 -- Dùng LEFT JOIN 2 SELECT p.id, p.name 3 FROM products p 4 LEFT JOIN order_items oi ON p.id = oi.product_id 5 WHERE oi.product_id IS NULL; 6 7 -- Dùng NOT EXISTS 8 SELECT p.id, p.name 9 FROM products p 10 WHERE NOT EXISTS (11 SELECT 1 12 FROM order_items oi 13 WHERE oi.product_id = p.id 14); 15 </pre>	

	<p>How do you delete all duplicate records but keep one unique record?</p> <p>To delete duplicate records but keep one unique row, you can use a CTE (Common Table Expression) with the ROW_NUMBER() window function. The idea is to assign a unique sequential number to duplicates within each group of identical values and then delete rows where this number is greater than 1. This ensures that one row is kept while all duplicates are removed.</p> <p>Explanation:</p> <ul style="list-style-type: none"> - PARTITION BY name, email groups duplicates based on these columns. - ROW_NUMBER() assigns 1 to the first row in each group. - DELETE removes rows where rn > 1, keeping only the first row. 	<pre>main.sql 1 ↴ WITH duplicates AS (2 ↴ SELECT id, 3 ↴ ROW_NUMBER() OVER (PARTITION BY name, email ORDER BY id) AS rn 4 ↴ FROM users 5 ↴) 6 ↴ DELETE FROM users 7 ↴ WHERE id IN (8 ↴ SELECT id 9 ↴ FROM duplicates 10 ↴ WHERE rn > 1 11); 12 </pre>
	<p>Write a query to calculate sales statistics by month and year.</p> <p>To calculate sales statistics by month and year, you can use the EXTRACT() function in PostgreSQL to extract the year and month from a date column, then group by these values. You can use aggregate functions like SUM() for total sales, COUNT() for the number of orders, and AVG() for average sales. This is useful for monthly or yearly reporting.</p> <p>Quick notes: EXTRACT(YEAR FROM ...) and EXTRACT(MONTH FROM ...) get year and month. Aggregate functions summarize the data. GROUP BY allows monthly/yearly aggregation.</p>	<pre>main.sql 1 SELECT 2 EXTRACT(YEAR FROM order_date) AS year, 3 EXTRACT(MONTH FROM order_date) AS month, 4 COUNT(*) AS total_orders, 5 SUM(total_amount) AS total_sales, 6 AVG(total_amount) AS avg_sales 7 FROM orders 8 GROUP BY year, month 9 ORDER BY year, month; 10 </pre>
	<p>Write a query to calculate the total, average, maximum, and minimum for each group.</p> <p>To calculate total, average, maximum, and minimum for each group, you use SQL aggregate functions (SUM(), AVG(), MAX(), MIN()) together with a GROUP BY clause. This allows you to summarize data per group, for example per department, category, or any grouping column.</p>	<pre>main.sql 1 SELECT 2 department_id, 3 SUM(salary) AS total_salary, 4 AVG(salary) AS avg_salary, 5 MAX(salary) AS max_salary, 6 MIN(salary) AS min_salary 7 FROM employee 8 GROUP BY department_id 9 ORDER BY department_id; 10 </pre>

	<p>How do you combine multiple columns into a single column?</p> <p>To combine multiple columns into a single column in SQL, you can use concatenation functions. In PostgreSQL, you can use the <code> </code> operator or the <code>CONCAT()</code> function. This is useful when you want to merge, for example, first name and last name into a full name column. You can also include separators like spaces or commas.</p> <p>Quick notes:</p> <ul style="list-style-type: none"> - <code>CONCAT()</code> = merges multiple columns/values - <code> </code> = PostgreSQL operator for string concatenation - Include separators (space, comma) as needed 	<pre>main.sql 1 -- Dùng hàm CONCAT 2 SELECT CONCAT(first_name, ' ', last_name) AS full_name 3 FROM employees; 4 5 -- Dùng toán tử 6 SELECT first_name ' ' last_name AS full_name 7 FROM employees; 8</pre>
	<p>Explain the differences between temporary tables, table variables, and CTEs (Common Table Expressions).</p> <p>In SQL, temporary tables, table variables, and CTEs (Common Table Expressions) are all ways to store intermediate results, but they differ in scope, lifetime, and usage.</p> <ul style="list-style-type: none"> - Temporary tables (<code>CREATE TEMP TABLE</code>) are stored in the database like regular tables but are automatically dropped at the end of the session or optionally at the end of a transaction. They can have indexes, constraints, and statistics, and can be referenced multiple times in queries. They are useful for large datasets or multiple-query processing. - Table variables (e.g., in T-SQL <code>DECLARE @table TABLE</code>) are variables that store a table structure in memory. They have limited scope (usually the batch or procedure) and generally do not maintain statistics or allow indexes except primary keys. They are best for small datasets and lightweight temporary storage. - CTEs (Common Table Expressions) are temporary result sets defined using <code>WITH</code>. They exist only for the duration of a single query and are not physically stored. They are mainly used for readability, recursion, or simplifying complex queries. 	<pre>main.sql 1 -- Temporary table 2 CREATE TEMP TABLE temp_orders AS 3 SELECT * FROM orders WHERE order_date >= '2025-01-01'; 4 5 -- CTE 6 WITH recent_orders AS (7 SELECT * FROM orders WHERE order_date >= '2025-01-01' 8) 9 SELECT customer_id, COUNT(*) 10 FROM recent_orders 11 GROUP BY customer_id; 12</pre>

	<p>How do you implement recursive queries in SQL?</p> <p>In SQL, you implement recursive queries using CTEs (Common Table Expressions) with the WITH RECURSIVE clause. Recursive queries are useful for hierarchical data, like organizational charts, bill-of-materials, or tree structures. A recursive CTE consists of two parts:</p> <ol style="list-style-type: none"> 1. Anchor member: the base query that provides the starting row (s). 2. Recursive member: references the CTE itself to repeatedly fetch child or related rows. <p>The recursion stops when no new rows are returned.</p>	<pre>main.sql 1 -- Giả sử bảng employees có các cột id và manager_id: 2 WITH RECURSIVE employee_hierarchy AS (3 -- Anchor member: chọn quản lý cấp cao nhất 4 SELECT id, name, manager_id, 1 AS level 5 FROM employees 6 WHERE manager_id IS NULL 7 8 UNION ALL 9 10 -- Recursive member: chọn nhân viên thuộc quản lý ở cấp trước 11 SELECT e.id, e.name, e.manager_id, eh.level + 1 12 FROM employees e 13 INNER JOIN employee_hierarchy eh ON e.manager_id = eh.id 14) 15 SELECT * 16 FROM employee_hierarchy 17 ORDER BY level, manager_id;</pre>
	<p>Windowing trong SQL là gì, và nó khác gì so với GROUP BY?</p> <p>Windowing trong SQL là các hàm cửa sổ (window functions), dùng để thực hiện tính toán trên một tập hợp các dòng liên quan đến dòng hiện tại, không gộp các dòng thành một kết quả duy nhất. Khác với GROUP BY, gộp tất cả dòng thành một kết quả cho mỗi nhóm, window functions giữ nguyên số dòng gốc đồng thời cung cấp tổng hợp, xếp hạng, hoặc trung bình động trên "cửa sổ" các dòng.</p> <p>Các hàm window phổ biến: ROW_NUMBER(), RANK(), SUM() OVER(...), AVG() OVER(...), LEAD()/LAG(). Câu lệnh OVER() xác định cửa sổ, có thể chia nhóm (PARTITION BY) và sắp xếp (ORDER BY).</p> <p>Lưu ý nhanh:</p> <p>GROUP BY = gộp các dòng thành một dòng mỗi nhóm Window function = tính toán không giảm số dòng Dùng cho xếp hạng, tổng lũy kế, trung bình động, v.v.</p>	<pre>main.sql 1 -- Tính tổng lương theo phòng ban nhưng giữ từng dòng riêng 2 SELECT 3 id, 4 name, 5 department_id, 6 salary, 7 SUM(salary) OVER (PARTITION BY department_id) AS total_department_salary, 8 ROW_NUMBER() OVER (PARTITION BY department_id ORDER BY salary DESC) AS rn 9 10 FROM employees;</pre>

	<p>Explain the difference between RANK() and NTILE().</p> <p>Both RANK() and NTILE() are window functions in SQL used to assign numbers to rows, but they behave differently.</p> <ul style="list-style-type: none"> - RANK() assigns a rank to each row based on the ordering of a column. Rows with the same value get the same rank, and gaps appear in the ranking sequence after ties. - NTILE(n) divides the ordered rows into n approximately equal groups (buckets) and assigns a bucket number to each row. It does not care about ties but ensures all rows are distributed into the specified number of tiles. <p>Quick notes:</p> <ul style="list-style-type: none"> RANK = ranking with possible gaps for ties NTILE = splits rows into equal-sized buckets Use RANK for positions, NTILE for grouping into percentiles/quartiles 	
	<p>How can you pivot rows into columns in SQL?</p> <p>To pivot rows into columns in SQL, you transform row values into column headers. In PostgreSQL, you can use the crosstab() function from the tablefunc extension or simulate it using CASE with GROUP BY. Pivoting is useful for reporting or summarizing data, for example turning monthly sales into separate columns for each month.</p>	<pre>main.sql 1 -- Ví dụ RANK 2 SELECT name, salary, 3 RANK() OVER (ORDER BY salary DESC) AS rk 4 FROM employees; 5 6 -- Ví dụ NTILE (chia thành 4 phần) 7 SELECT name, salary, 8 NTILE(4) OVER (ORDER BY salary DESC) AS quartile 9 FROM employees; 10 </pre>
	<p>How can you unpivot columns into rows in SQL?</p> <p>To unpivot columns into rows in SQL, you transform multiple column values into a single column with multiple rows. In PostgreSQL, you can use the UNION ALL approach or the jsonb/lateral methods. Unpivoting is useful when you need to normalize data for analysis, like converting monthly sales columns into a single column with month and sales values.</p>	<pre>main.sql 1 -- Ví dụ PostgreSQL dùng CASE 2 SELECT product_id, 3 SUM(CASE WHEN EXTRACT(MONTH FROM order_date) = 1 THEN total_amount ELSE 0 END) AS 4 jan_sales, 5 SUM(CASE WHEN EXTRACT(MONTH FROM order_date) = 2 THEN total_amount ELSE 0 END) AS 6 feb_sales, 7 SUM(CASE WHEN EXTRACT(MONTH FROM order_date) = 3 THEN total_amount ELSE 0 END) AS 8 mar_sales 9 FROM orders 10 WHERE product_id = 1 11 ORDER BY product_id; 12 13 -- Ví dụ PostgreSQL dùng crosstab() 14 -- Kích hoạt extension tablefunc 15 CREATE EXTENSION IF NOT EXISTS tablefunc; 16 SELECT * FROM crosstab(17 'SELECT product_id, EXTRACT(MONTH FROM order_date) AS month, SUM(total_amount) 18 FROM orders 19 GROUP BY product_id, month 20 ORDER BY product_id, month' 21) AS ct (product_id INT, jan NUMERIC, feb NUMERIC, mar NUMERIC); 22 </pre> <pre>main.sql 1 -- Ví dụ PostgreSQL dùng UNION ALL 2 SELECT product_id, 'jan' AS month, jan_sales AS sales FROM product_sales 3 UNION ALL 4 SELECT product_id, 'feb' AS month, feb_sales AS sales FROM product_sales 5 UNION ALL 6 SELECT product_id, 'mar' AS month, mar_sales AS sales FROM product_sales; 7 8 -- Ví dụ PostgreSQL dùng VALUES + LATERAL 9 SELECT ps.product_id, v.month, v.sales 10 FROM product_sales ps, 11 LATERAL (12 VALUES 13 ('jan', jan_sales), 14 ('feb', feb_sales), 15 ('mar', mar_sales) 16) AS v(month, sales); 17 </pre>

	<p>Explain the differences between INNER JOIN with multiple tables versus multiple subqueries.</p> <p>When retrieving data from multiple tables, you can either use INNER JOINS or multiple subqueries, but there are key differences.</p> <ul style="list-style-type: none"> - INNER JOIN with multiple tables combines rows from all tables based on matching keys in a single query. It is usually more efficient, easier to read, and allows the query planner to optimize joins effectively. It returns a single result set with columns from all joined tables. - Multiple subqueries involve executing separate queries, often in WHERE clauses or in the SELECT statement. This can be less efficient, especially if the subqueries run repeatedly for each row, and can be harder to read. Subqueries are useful when you need aggregated or filtered results independently before joining them to the main query. 	<pre>main.sql 1 -- INNER JOIN nhiều bảng 2 SELECT o.id AS order_id, c.name AS customer_name, p.name AS product_name 3 FROM orders o 4 INNER JOIN customers c ON o.customer_id = c.id 5 INNER JOIN order_items oi ON o.id = oi.order_id 6 INNER JOIN products p ON oi.product_id = p.id; 7 8 -- Dùng subquery 9 SELECT o.id AS order_id, 10 (SELECT name FROM customers WHERE id = o.customer_id) AS customer_name, 11 (SELECT name FROM products WHERE id = oi.product_id) AS product_name 12 FROM orders o 13 INNER JOIN order_items oi ON o.id = oi.order_id; 14</pre>
	<p>How do you handle hierarchical or tree-structured data in SQL?</p> <p>Hierarchical or tree-structured data in SQL, like organizational charts or categories, can be handled using recursive queries with CTEs (Common Table Expressions). The idea is to start with root nodes (top-level items) as the anchor, then recursively retrieve child nodes until all levels are fetched. You can also use parent_id/child_id relationships and window functions for path or level calculations.</p>	<pre>main.sql 1 -- Giả sử bảng categories có id và parent_id: 2 WITH RECURSIVE category_tree AS (3 -- Anchor: danh mục cấp cao nhất 4 SELECT id, name, parent_id, 1 AS level, name AS path 5 FROM categories 6 WHERE parent_id IS NULL 7 8 UNION ALL 9 10 -- Recursive: các mục con 11 SELECT c.id, c.name, c.parent_id, ct.level + 1, ct.path ' > ' c.name 12 FROM categories c 13 INNER JOIN category_tree ct ON c.parent_id = ct.id 14) 15 SELECT * 16 FROM category_tree 17 ORDER BY path; 18 </pre>
	<p>How can you implement soft delete (logical delete) in a table?</p> <p>A soft delete (logical delete) means you mark a row as deleted instead of physically removing it from the table. This is usually done by adding a column like is_deleted or deleted_at. Queries then filter out deleted rows by checking this column. Soft delete is useful for auditing, recovery, or preserving historical data.</p> <p>Alternative approach: use a timestamp column deleted_at instead of a boolean to record when a row was deleted.</p> <p>Quick notes:</p> <ul style="list-style-type: none"> Soft delete = mark row, do not physically delete Filter queries to exclude deleted rows Allows recovery and auditing 	<pre>main.sql 1 -- Thêm cột đánh dấu dòng bị xóa 2 ALTER TABLE employees ADD COLUMN is_deleted BOOLEAN DEFAULT FALSE; 3 4 -- Soft delete một dòng 5 UPDATE employees 6 SET is_deleted = TRUE 7 WHERE id = 123; 8 9 -- Truy vấn các dòng còn hiệu lực (chưa xóa) 10 SELECT * 11 FROM employees 12 WHERE is_deleted = FALSE; 13</pre>

	<p>How can you identify and optimize queries with high I/O usage?</p> <p>Queries with high I/O usage can slow down your database because they read or write a large amount of data from disk. To identify them, you can use execution plans, database statistics, or monitoring tools. Common signs include sequential scans on large tables, large joins, or frequent table scans instead of index usage.</p> <p>To optimize high I/O queries:</p> <ol style="list-style-type: none"> 1. Use indexes on columns frequently used in WHERE, JOIN, or ORDER BY clauses. 2. Limit scanned data with proper filtering conditions. 3. Avoid unnecessary SELECT *; fetch only needed columns. 4. Break complex queries into smaller steps or pre-aggregate data if possible. 5. Analyze query plans (EXPLAIN) to see which operations cause high I/O. <p>Quick notes:</p> <ul style="list-style-type: none"> Look for sequential scans on large tables Check the estimated rows vs actual rows in EXPLAIN Adding indexes or rewriting queries can reduce I/O 	<pre><code>main.sql 1 -- Phân tích kế hoạch truy vấn 2 EXPLAIN ANALYZE 3 SELECT * 4 FROM orders o 5 JOIN customers c ON o.customer_id = c.id 6 WHERE o.order_date >= '2025-01-01'; 7 </code></pre>
	<p>Explain the difference between NOLOCK (or READ UNCOMMITTED) and default isolation levels.</p> <p>NOLOCK (or READ UNCOMMITTED) is a transaction isolation level that allows a query to read data without acquiring shared locks. This means it can read uncommitted or "dirty" data from other transactions. It improves performance and reduces blocking but can result in inaccurate or inconsistent results.</p> <p>The default isolation level in most databases (e.g., READ COMMITTED in PostgreSQL, SQL Server) ensures that a query only reads committed data, preventing dirty reads. This is safer but can cause locks and blocking if many transactions are running concurrently.</p> <p>Quick notes:</p> <ul style="list-style-type: none"> NOLOCK = faster, non-blocking, may read dirty data Default isolation = safer, reads only committed data, may block Use NOLOCK carefully; avoid for critical financial or audit data 	<pre><code>main.sql 1 -- Dùng NOLOCK / READ UNCOMMITTED 2 SELECT * 3 FROM orders WITH (NOLOCK) 4 WHERE order_date >= '2025-01-01'; 5 6 -- Mức READ COMMITTED mặc định (không cần hint) 7 SELECT * 8 FROM orders 9 WHERE order_date >= '2025-01-01'; 10 </code></pre>

	<p>How do you merge data from two tables (UPSERT) in SQL?</p> <p>UPSERT in SQL is a combination of INSERT and UPDATE: it inserts a new row if it doesn't exist, or updates the existing row if it does. In PostgreSQL, this can be achieved using INSERT ... ON CONFLICT, specifying the unique constraint or primary key to detect conflicts. UPSERT is useful for synchronizing or merging data between tables.</p> <p>Quick notes:</p> <ul style="list-style-type: none"> - UPSERT = insert if new, update if exists - ON CONFLICT (key) detects duplicate keys - EXCLUDED refers to the row that would have been inserted 	<pre>main.sql 1 -- Giả sử bạn có hai bảng: source_table và target_table với primary key id. 2 INSERT INTO target_table (id, name, value) 3 SELECT id, name, value 4 FROM source_table 5 ON CONFLICT (id) 6 DO UPDATE SET 7 name = EXCLUDED.name, 8 value = EXCLUDED.value; 9</pre>
	<p>Explain how indexing affects INSERT, update, and delete performance.</p> <p>Indexing improves SELECT performance by allowing the database to quickly locate rows, but it also affects INSERT, update, and delete operations:</p> <ul style="list-style-type: none"> - INSERT: When inserting a new row, all relevant indexes must also be updated. More indexes = more overhead, which can slow down inserts. - UPDATE: Updating indexed columns requires updating the index as well. If many indexes exist on the updated column(s), the operation becomes slower. - DELETE: Deleting a row requires removing entries from all indexes that reference that row, which adds overhead. <p>In summary, indexes improve read performance but add extra work for write operations. It's important to balance between query speed and write performance.</p> <p>Quick notes:</p> <ul style="list-style-type: none"> - More indexes = faster reads, slower writes - Only index columns used frequently in queries - Consider partial or filtered indexes if appropriate 	
	<p>Question: What is the result of NULL = NULL?</p> <p>Trick: Nhiều người nghĩ là TRUE. Thực tế: kết quả là NULL (UNKNOWN). So sánh phải dùng IS NULL.</p>	

	<p>Question: What is the difference between COUNT(column_name) and COUNT(*)?</p> <p>Trick: COUNT(column_name) chỉ đếm các giá trị không NULL, COUNT(*) đếm tất cả dòng.</p>	
	<p>Question: Can you use HAVING without GROUP BY?</p> <p>Trick: Có thể, nhưng hiếm. Nhiều người sẽ trả lời là không ngay lập tức. SQL cho phép HAVING mà không có GROUP BY (tương đương một nhóm ảo).</p>	
	<p>Question: What happens if a subquery in WHERE returns multiple rows with = operator?</p> <p>Trick: Nó sẽ lỗi. Để tránh phải dùng IN thay vì =.</p>	
	<p>Question: Can you select a column in SQL that is not part of GROUP BY and not an aggregate?</p> <p>Trick: Một số hệ quản trị (Oracle strict mode) sẽ lỗi; MySQL cho phép nhưng kết quả không xác định.</p>	
	<p>Question: What is the difference between INNER JOIN ON 1=1 and CROSS JOIN?</p> <p>Trick: Nhiều người không nhận ra rằng INNER JOIN ON 1=1 thực chất tạo ra Cartesian product như CROSS JOIN.</p>	
	<p>Question: What is the purpose of ORDER BY NULL?</p> <p>Trick: Nó tắt việc sắp xếp trong một query, thường dùng để tối ưu aggregate queries.</p>	
	<p>Question: What happens if you UPDATE a table using a JOIN that matches multiple rows per record?</p> <p>Trick: Kết quả có thể bị cập nhật nhiều lần hoặc không xác định, tùy DBMS.</p>	
	<p>Question: If two transactions update the same row concurrently, which value will persist?</p> <p>Trick: Phụ thuộc vào isolation level; nhiều người trả lời trực quan nhưng quên rằng isolation level quyết định lock/rollback.</p>	
	<p>Question: What happens if two tables in a JOIN have the same column name and you don't alias?</p> <p>Trick: Một số DBMS lỗi ngay, một số trả về không xác định.</p>	

	<p>Question: What is the result of '5' + 1 in SQL?</p> <p>Trick: Tùy DBMS: MySQL auto convert, SQL Server lỗi hoặc convert khác. Rất nhiều người quên về kiểu dữ liệu.</p>	
	<p>Question: Which is faster: EXISTS or IN?</p> <p>Trick: Không có câu trả lời chung; phụ thuộc DBMS, số lượng bản ghi, và NULL trong danh sách. Nhiều người trả lời vội vàng là EXISTS luôn nhanh hơn.</p>	
	<p>Question: What is the result of this query?</p> <pre>SELECT COUNT(*) FROM employees WHERE salary > 5000 OR salary < 5000;</pre> <p>Trick: Nhiều người nghĩ là toàn bộ nhân viên, nhưng nếu có salary IS NULL, họ không được tính. Vì NULL không so sánh được bằng toán tử > hoặc <.</p>	
	<p>Question: What does this return?</p> <pre>SELECT * FROM employees WHERE NOT (salary > 5000);</pre> <p>Trick: Không tương đương với salary <= 5000 vì các hàng có NULL cũng bị loại bỏ.</p>	
	<p>Question: What is the output of:</p> <pre>SELECT 1 / 0;</pre> <p>Trick: Một số DBMS sẽ báo lỗi (division by zero), nhưng MySQL trong chế độ mặc định có thể trả NULL.</p>	
	<p>Question: Can you use GROUP BY without any aggregate function?</p> <p>Answer: Có thể. Nó chỉ nhóm các hàng mà không tính toán, thường dùng để loại bỏ trùng (như DISTINCT).</p>	
	<p>Question: What does SELECT DISTINCT NULL, NULL; return?</p> <p>Trick: Chỉ 1 dòng — DISTINCT coi tất cả NULL là "bằng nhau".</p>	
	<p>Question: Which returns more rows, INNER JOIN or LEFT JOIN?</p> <p>Trick: Không có câu trả lời tuyệt đối — LEFT JOIN có thể nhiều hơn, nhưng nếu mọi hàng có match, thì số dòng bằng nhau.</p>	

	<p>Question: Does HAVING filter before or after aggregation?</p> <p>Trick: Sau khi nhóm (GROUP BY) và tính toán aggregate xong mới lọc.</p>	
	<p>Question: What happens when you compare NULL with anything using !=?</p> <p>Trick: Luôn trả UNKNOWN. NULL không bao giờ bằng hay khác bất kỳ giá trị nào.</p>	
	<p>Question: Does ORDER BY guarantee consistent ordering if there are duplicate values in the sorting column?</p> <p>Trick: Không. Thứ tự các hàng có cùng giá trị có thể khác nhau giữa các lần chạy nếu không có tiêu chí phụ.</p>	
	<p>Question: Is DELETE FROM table; the same as TRUNCATE TABLE table;?</p> <p>Trick: Không. DELETE có thể rollback và kích hoạt trigger; TRUNCATE thì không (tùy DBMS).</p>	
	<p>Question: Can NULL be a value in a UNIQUE column?</p> <p>Trick: Có thể. SQL coi NULL là “unknown”, nên hai NULL không bị coi là trùng (MySQL, PostgreSQL).</p>	
	<p>Question: Which one executes faster — SELECT * or SELECT specific_columns?</p> <p>Trick: SELECT * có thể chậm hơn vì phải fetch tất cả các cột, đặc biệt với nhiều cột hoặc I/O lớn.</p>	
	<p>Question: What will this return?</p> <pre>SELECT SUM(CASE WHEN dept_id = 10 THEN 1 END) AS count_dept10 FROM employees;</pre> <p>Trick: Có thể trả NULL nếu không có dòng nào thỏa dept_id = 10, vì SUM(NULL) là NULL, không phải 0.</p>	
	<p>Question: Can you use an aggregate function inside a WHERE clause?</p> <p>Trick: Không. WHERE chạy trước aggregation. Phải dùng HAVING hoặc subquery.</p>	

	<p>Question: What is the difference between:</p> <pre>SELECT COUNT(*) FROM table; SELECT COUNT(1) FROM table;</pre> <p>Trick: Thường không khác biệt về hiệu năng. Một số người tin COUNT(1) nhanh hơn, nhưng hầu hết DBMS tối ưu như nhau.</p>	
	<p>Question: Will this query run?</p> <pre>SELECT * FROM employees WHERE salary = ALL (SELECT salary FROM employees);</pre> <p>Trick: Có thể trả rỗng vì điều kiện nghĩa là “salary bằng tất cả các salary khác” — chỉ đúng nếu mọi người cùng lương.</p>	
	<p>Question: What happens if you update a table that is used in its own subquery?</p> <pre>UPDATE employees SET salary = salary * 1.1 WHERE emp_id IN (SELECT manager_id FROM employees);</pre> <p>Trick: Một số DBMS (như MySQL) không cho phép subquery đọc và ghi cùng bảng (error 1093).</p>	
	<p>Question: What's wrong with this query?</p> <pre>SELECT name, AVG(salary) FROM employees;</pre> <p>Trick: Không có GROUP BY name → lỗi (trừ khi MySQL trong mode lỏng).</p>	
	<p>Question: What's the result of SELECT NULL + 5;?</p> <p>Trick: NULL. Mọi phép toán với NULL → NULL.</p>	
	<p>Question: What happens when you insert a row with no value into an IDENTITY (auto-increment) column?</p> <p>Trick: DBMS tự tạo giá trị mới, nhưng nếu bạn chèn NULL hoặc DEFAULT, hành vi có thể khác tùy hệ quản trị.</p>	